



METIS

---

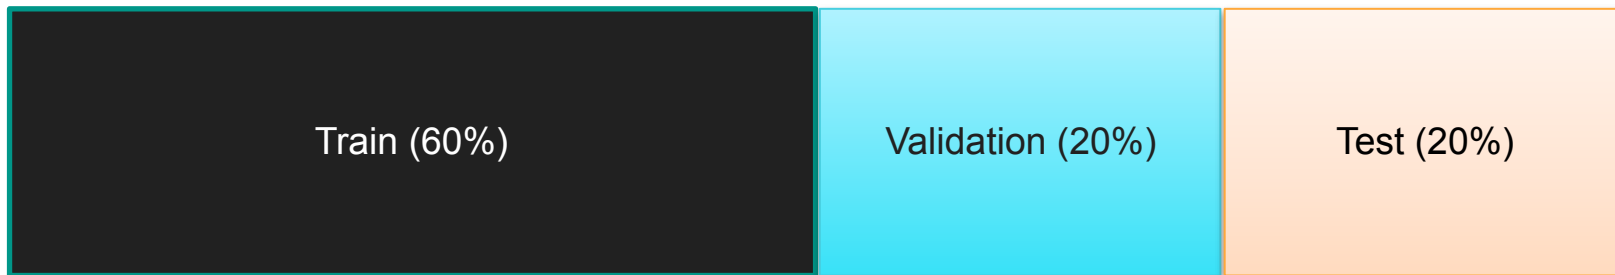
# Validation And Testing: **Selecting Models And Estimating Their Quality**

---

# LEARNING GOALS



1. Understand the importance of estimating model generalization error (testing) and practical methodologies for doing so
2. Learn validation/cross-validation strategies for model selection
3. Recognize the difference between validation and testing



---

# Testing: Estimating Generalization Error

---

METIS



**Generalization Error:** How well can we expect a model to perform on new data from the same distribution as the training data?

- Predictive models are only *useful* if they can give us good target approximations for samples that we haven't seen before
- Example: Zillow predicts the market value of a home before it's listed for sale, training a model on known listing prices
- So when evaluating models, we should attempt to measure how well they *generalize*, i.e. estimate performance on samples we didn't train on. We call this **testing**.

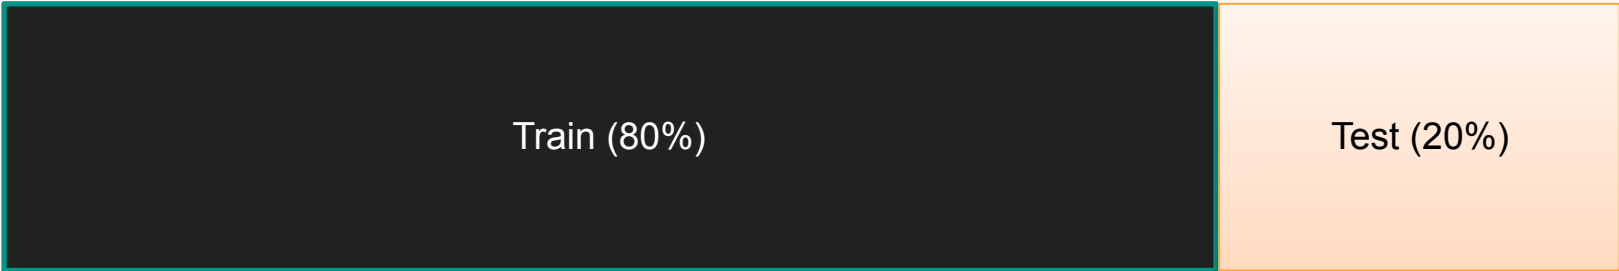
# TESTING, IN PRACTICE



**Simulate generalization:** We can *hold out* a portion of our labeled dataset to simulate the real-world challenge of unseen samples

We call this a **test set**, and exclude it from the data we train on

We then *estimate generalization error* by making predictions on test, and scoring those predictions against the ground truth (our test labels)



Train (80%)

Test (20%)

# TESTING, IN PRACTICE; cont.



## 1. Fit model to training data



Train (80%)

Test (20%)

# TESTING, IN PRACTICE; cont.



1. Fit model to training data
2. Score model on testing data



Train (80%)


Test (20%)

# TEST USUALLY UNDERPERFORMS TRAIN



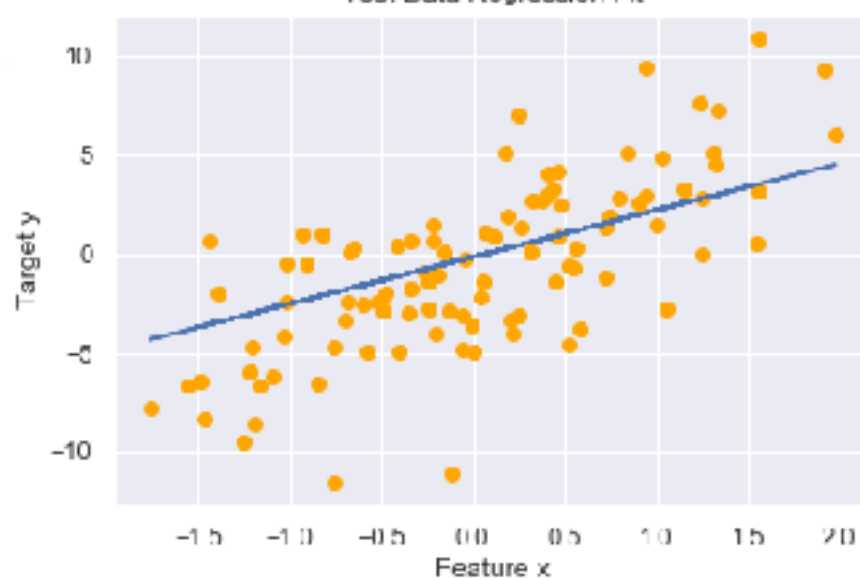
**Model is optimized to perform as well as possible on train**, so it's no surprise that it tends to have a worse evaluation score on test (though this is not guaranteed).

Train Data Regression Fit



→  $R^2$ : .49

Test Data Regression Fit




→  $R^2$ : .43




# THE RISK OF OVERFITTING



**Out of sample evaluations let us check for overfitting:** more complex models can get arbitrarily better at predicting the train data, but will start to fit to spurious patterns and generalize more poorly



→ Degree 4  $R^2$ : .54



→ Degree 4  $R^2$ : .42

---

# Validation: Optimizing Our Modeling Choices

---

METIS

# VALIDATION: OPTIMIZING CHOICES



When we construct predictive models, we typically have **many choices**:

- Features: which data columns do we include/exclude or engineer?
- Preprocessing: how should we handle nulls? Should we standardized the scale of the features?
- Hyper-parameters: What degree polynomial regression should we fit? What regularization strength should we use? How does a random forest model compare to a linear regression model?

# VALIDATION IN PRACTICE



We can make some choices using our domain knowledge and good instincts, but a **validation framework** gives us an empirical way to choose and avoid over/under-fitting

We validate with the usual *best generalization* end-goal in mind: we exclude validation data from training, and use it to score predictions across a range of model choices

We can then *select* a choice of model based on the strongest validation score - i.e., this score gives us **direct feedback on a possible choice**. Once we've chosen a model, we can combine our train and validation sets, retrain the model, and get the test score

Train (60%)

Validation (20%)

Test (20%)

# VALIDATION IN PRACTICE; cont.



## 1. Train candidate models

- Linear Regression
- Polynomial Regression
- Ridge Regression



Train (60%)

Validation (20%)

Test (20%)

# VALIDATION IN PRACTICE; cont.



## 1. Train candidate models

- Linear Regression
- Polynomial Regression
- Ridge Regression

## 2. Score candidates

→  $R^2$ : .35

→  $R^2$ : .50

→  $R^2$ : .25

Train (60%)

Validation (20%)

Test (20%)



# VALIDATION IN PRACTICE; cont.



## 3. Retrain best candidate on train + validation

→ Polynomial Regression



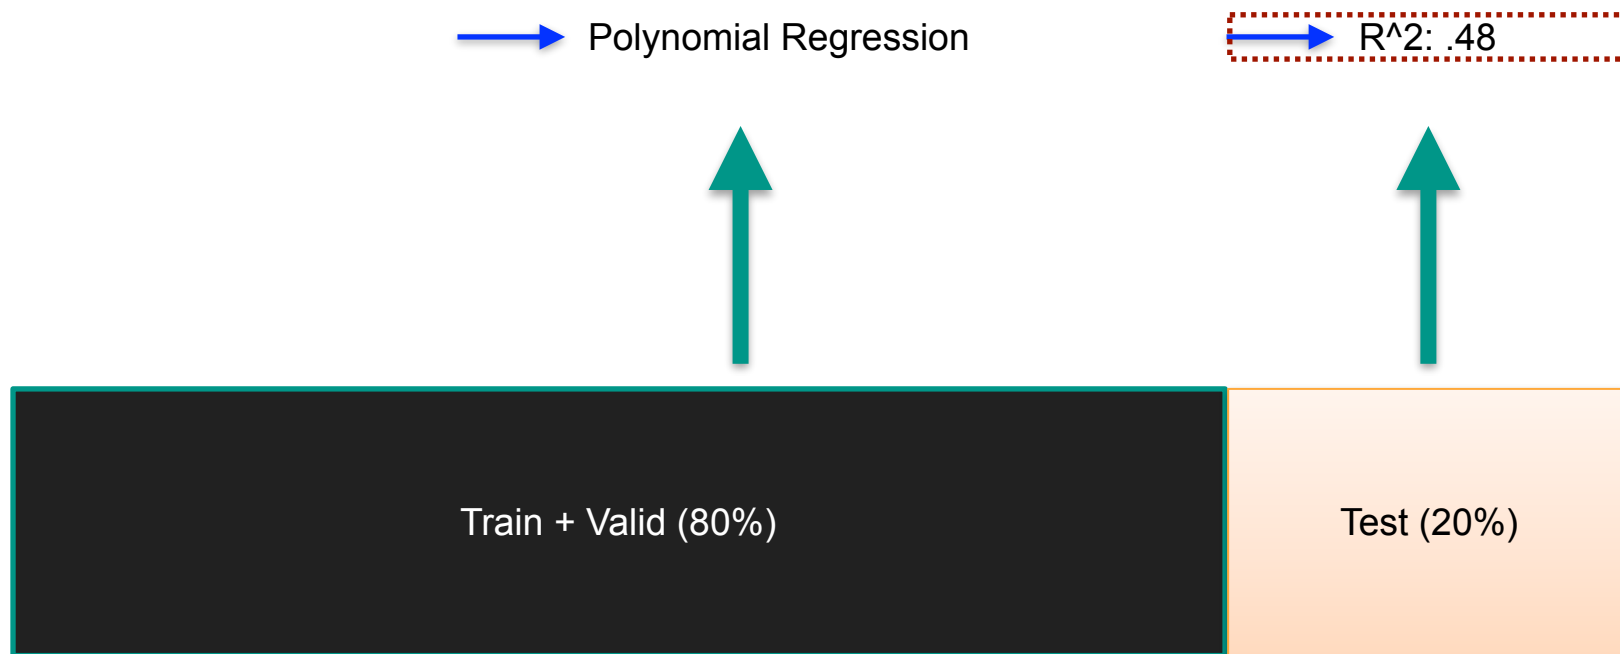
Train + Valid (80%)

Test (20%)

# VALIDATION IN PRACTICE; cont.



## 4. Score final model on test





# VALIDATION: KEY CONSIDERATIONS



**Validation is not testing:** This is a very common pitfall. Once you've used a data set to influence your model choices through direct feedback, it can't be used to give an unbiased estimate of generalization error

**Fair comparisons:** Candidate models should be compared against the same validation scheme

**Data efficiency:** after we use a portion of our data for validation, we should reuse it as training data to improve the final model

---

# Cross Validation: Optimizing Our Optimization Of Choices

---

METIS

# CROSS VALIDATION: ADDING MORE RIGOR



**Cross-validation is about reliability and efficiency:** what if we overfit to an unlucky validation set? Can we use more of our data than just one hold-out for validation?

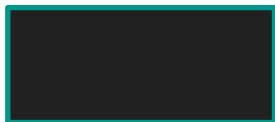
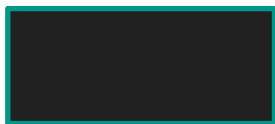
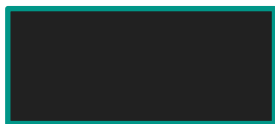
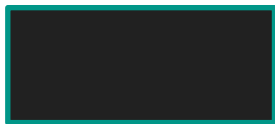
**K-Fold partitioning:** randomly divide our non-test data into  $K$  equal-sized groups. Each group will be used as a validation set once, and we'll compare candidate models via mean scores across all validation scores.

**K is usually 5 or 10:** depends on problem and size of data, but these are common choices

# CROSS VALIDATION IN PRACTICE



Train data:  
divided into  
5 folds



Do the following,  
for each  
candidate model

Test, held out

# CROSS VALIDATION IN PRACTICE

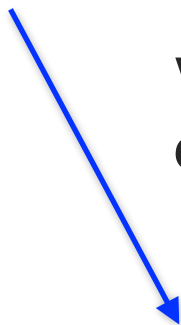


Train data:  
divided into  
5 folds



**Train** on  
black folds

**Validate** on  
cyan fold

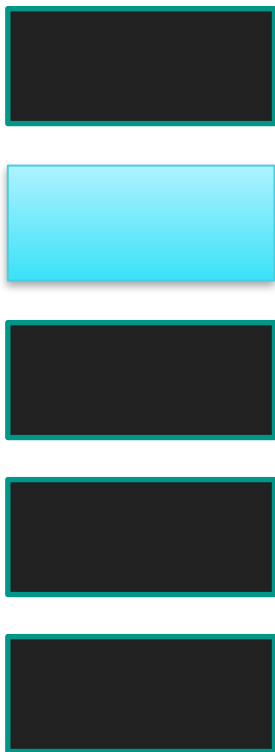


Validation  $R^2$  scores: .45,

# CROSS VALIDATION IN PRACTICE




Train data:  
divided into  
5 folds



**Train** on  
black folds

**Validate** on  
cyan fold

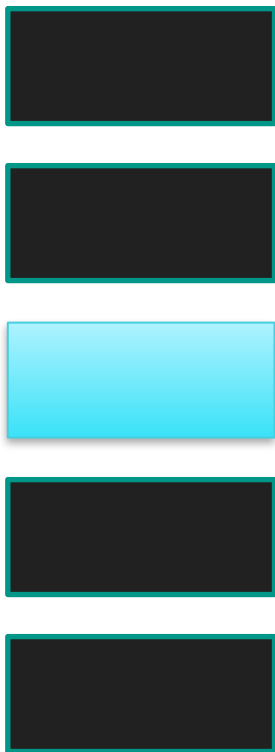


**Validation  $R^2$  scores:** .45, .39,

# CROSS VALIDATION IN PRACTICE




Train data:  
divided into  
5 folds



**Train** on  
black folds

**Validate** on  
cyan fold

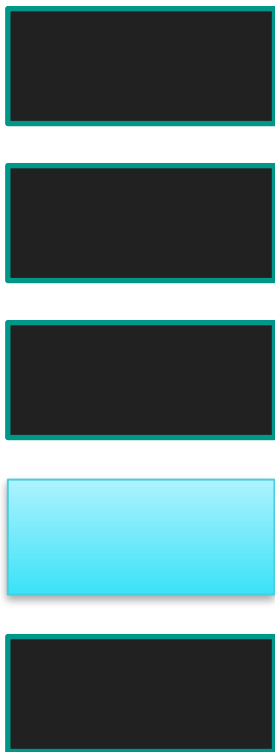


**Validation  $R^2$  scores:** .45, .39, .57,

# CROSS VALIDATION IN PRACTICE

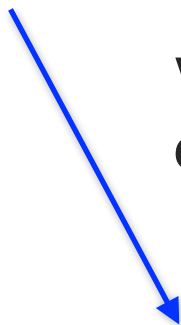


Train data:  
divided into  
5 folds



**Train** on  
black folds

**Validate** on  
cyan fold



**Validation  $R^2$  scores:** .45, .39, .57, .49,



# CROSS VALIDATION IN PRACTICE

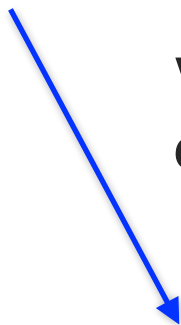


Train data:  
divided into  
5 folds



**Train** on  
black folds

**Validate** on  
cyan fold

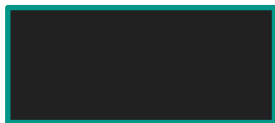
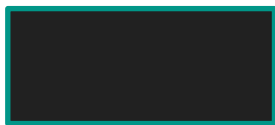

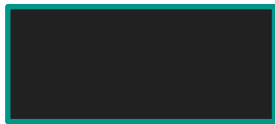
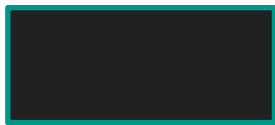


**Validation  $R^2$  scores:** .45, .39, .57, .49, .50

# CROSS VALIDATION IN PRACTICE



Train data:  
divided into  
5 folds



Produces a set of results for each  
candidate model

Linear regression



Validation  $R^2$  scores: .45, .39, .57, .49, .50

Poly regression

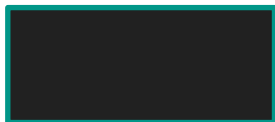
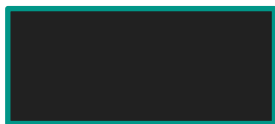
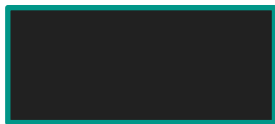


Validation  $R^2$  scores: .53, .43, .67, .55, .51

# CROSS VALIDATION IN PRACTICE



Train data:  
divided into  
5 folds



Summarize candidates by mean  
score, select best

Linear regression



5-Fold validation mean  $R^2$  score: .48

Poly regression




5-Fold validation mean  $R^2$  score: .54

# CROSS VALIDATION IN PRACTICE



Train data:  
recombined



Polynomial regression  
selected as best  
candidate model

**Poly regression,  
retrained on all data,  
final score on test**



Test, held out

$R^2$ : .48

---

# Validation And Testing: Recap

---

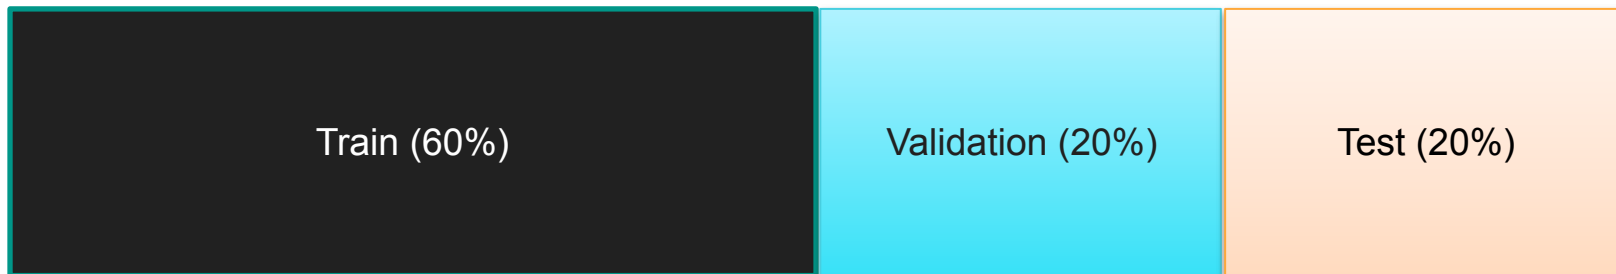
METIS

# WORKFLOW METHOD 1: Train/Valid/Test



Collect set of candidate models. Fit each on train, score on validation, select final model via best validation score

Retrain final model on train + validation, report score on test as estimate of generalization error



## WORKFLOW METHOD 2: CV/Test



Collect set of candidate models. Run each through a K-fold CV loop, select final model via best mean validation score

Retrain final model on combined folds, report score on test as estimate of generalization error



# VALIDATION VS. CV - WHEN TO USE?



**Simple validation** is significantly faster and often representative enough when working with very large samples (~millions+)

**Cross validation** is more appropriate with small-medium size data or when variance in results between different validation sets is high



# CAN WE PUSH EVEN FURTHER?



There's nothing stopping us from doing **repeated rounds of CV** with different random K-folds for even more rigor

- Also an alternate form of testing: run CV on all data to select, then run another K-Fold loop to get multiple/mean out of sample scores

We usually take means across validation folds in CV to compare model candidates, but we can also gain information from **distributions of scores across folds** (e.g. variance)

**Nested CV** is another advanced technique for a more robust combination of CV and testing: [see here for detail](#).

# WHAT ABOUT TIME SERIES DATA?



For certain data (e.g. data with a time series component), a more specialized validation and testing setup may be called for

In **time-based validation**, validation/testing data should chronologically follow training data: this simulates the model's **generalization into the future**

Example: when using historical stock market data to predict future returns, try to make sure that past patterns continue to be predictive in the future

Train (2015-2017)

Validation (2018)

Test (2019)

# SUMMARY



## Training

In sample

Model building

Optimize model  
parameters (fit)

## Validation

Out of sample

Feedback to model  
selection

Optimize choices:  
features and model  
hyper-parameters

## Testing

Out of sample

No feedback to model  
selection

Final estimate of model  
generalization error

A close-up photograph of a Shiba Inu dog's face, looking slightly to the left with wide, dark eyes. The dog has light tan fur with darker markings around its eyes and ears. The background is a soft, out-of-focus yellow. Three lines of white, bold, sans-serif text with black outlines are overlaid on the image. A small watermark is visible in the bottom left corner.

**VALIDATION SO WOW**

**SUCH MODEL SELECTION**

**VERY NOT ESTIMATE  
OF GENERALIZATION ERROR**