


# Stochastic Gradient Descent




# Gradient Descent

Start with a cost function  $J(\beta)$ :




# Gradient Descent

$$J(\beta_0, \beta_1)$$




# Gradient Descent

Then gradually move to the minimum.



# Gradient Descent with Linear Regression

How can we do this?






# Gradient Descent with Linear Regression

How can we do this?  
(without seeing the graph of  $J(\beta)$ !)

Start with the function  $J(\beta)$ :




# Gradient Descent with Linear Regression

How can we do this?

(without seeing the graph of  $J(\beta)$ !)

Start with the function  $J(\beta)$ :

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Gradient Descent with Linear Regression


How can we do this?

(without seeing the graph of  $J(\beta)$ !)

Start with the function  $J(\beta)$ :

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

and compute its  
gradient vector  $\nabla J(\beta)$ .





# Gradient Descent with Linear Regression

How can we do this?


(without seeing the graph of  $J(\beta)$ !)

Start with the function  $J(\beta)$ :

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$


and compute its  
gradient vector  $\nabla J(\beta)$ .

The gradient points  
in the “**direction of  
maximum increase**” of  $J$ .



# Gradient Descent with Linear Regression


$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Gradient Descent with Linear Regression

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$


$$w_1 = w_0 - \alpha \nabla 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Gradient Descent with Linear Regression

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$


$$w_1 = w_0 - \alpha \left( \frac{\partial}{\partial \beta_0}, \dots, \frac{\partial}{\partial \beta_n} \right) 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Gradient Descent with Linear Regression

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

$$w_1 = w_0 - \alpha \nabla 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$






# Gradient Descent with Linear Regression

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$


$$w_2 = w_1 - \alpha \nabla 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Gradient Descent with Linear Regression

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$


$$w_3 = w_2 - \alpha \nabla 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$




# Gradient Descent with Linear Regression

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$


$$w_4 = w_3 - \alpha \nabla 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Gradient Descent



# Gradient Descent






# Gradient Descent


$$h_{\theta}(x)$$

(for fixed  $\theta_0, \theta_1$ , this is a function of  $x$ )




$$J(\theta_0, \theta_1)$$

(function of the parameters  $\theta_0, \theta_1$ )



# Stochastic Gradient Descent


$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$



# Stochastic Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Instead of using all points to find the gradient,  
Only use a SINGLE point each time




# Stochastic Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Instead of using all points to find the gradient,  
Only use a SINGLE point each time

$$w_1 = w_0 - \alpha \nabla 1/2 \left( (\beta_0 + \beta_1 x_{obs}^{(0)}) - y_{obs}^{(0)} \right)^2$$




# Stochastic Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Instead of using all points to find the gradient,  
Only use a SINGLE point each time

$$w_2 = w_1 - \alpha \nabla 1/2 \left( (\beta_0 + \beta_1 x_{obs}^{(1)}) - y_{obs}^{(1)} \right)^2$$






## Stochastic Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Instead of using all points to find the gradient,  
Only use a SINGLE point each time

$$w_3 = w_2 - \alpha \nabla 1/2 \left( (\beta_0 + \beta_1 x_{obs}^{(2)}) - y_{obs}^{(2)} \right)^2$$




## Stochastic Gradient Descent

$$J(\beta_0, \beta_1) = \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Instead of using all points to find the gradient,  
Only use a SINGLE point each time

$$w_4 = w_3 - \alpha \nabla 1/2 \left( (\beta_0 + \beta_1 x_{obs}^{(3)}) - y_{obs}^{(3)} \right)^2$$



## Faster

Derivative of single point at each step (instead of 100K)

## Online Training

Only need to keep single point in memory

No need to store 100K rows, large data no problem

## Covers Many Algorithms

Gradient Descent is the bottleneck for linear algorithms

Can do Linear Regression, Logistic Regression, SVMs

# **Some Implementations**

# Some Implementations

```
from sklearn.linear_model import SGDRegressor
```


```
from sklearn.linear_model import SGDClassifier
```



```
from sklearn.linear_model import SGDRegressor
```

```
from sklearn.linear_model import SGDRegressor
```

```
SGDRegressor(loss='squared_loss')
```



Sum of squared errors


squared loss ==  
Linear Regression

Optimization  
Function:

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

```
from sklearn.linear_model import SGDRegressor
```

```
SGDRegressor(loss='epsilon_insensitive')
```



- Minimize:

$$\frac{1}{2} \|w\|^2 + C \sum_{i=1}^N (\xi_i + \xi_i^*)$$

- Constraints:

$$y_i - wx_i - b \leq \varepsilon + \xi_i$$

$$wx_i + b - y_i \leq \varepsilon + \xi_i^*$$

$$\xi_i, \xi_i^* \geq 0$$

Source: <http://kernelsvm.tripod.com>

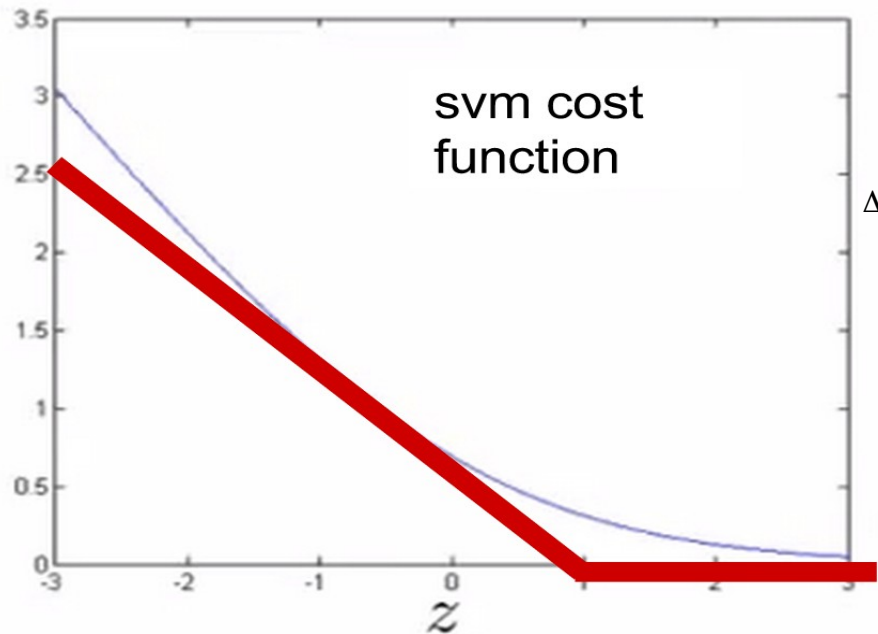
```
from sklearn.linear_model import SGDClassifier
```

```
from sklearn.linear_model import SGDClassifier
```

```
SGDClassifier(loss='hinge')
```

```
from sklearn.linear_model import SGDClassifier
```

```
SGDClassifier(loss='hinge')
```



Looks like a hinge.

hinge loss == SVM

Optimization  
Function:


$$Li = \sum_{j \neq y_i} \max(0, ((\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} + \Delta))$$

What is  $\Delta$  ? Just like with  $\lambda$  within our regularization term,  $\Delta$  it affects the trade-off between our data loss and our regularization loss within our objective function.



```
from sklearn.linear_model import SGDClassifier
```

```
SGDClassifier(loss='log')
```



This one's kind of clear

log loss == Logistic Regression

where

Optimization  
Function:

$$Li = -\log\left(\frac{e^{f_{yi}}}{\sum_j e^{f_j}}\right)$$

$$e^{f_{yi}} = \exp(\beta_0 + \omega^T x_{obs}^{(i)})$$

```
from sklearn.linear_model import SGDClassifier
```

```
SGDClassifier(alpha=0.0001,  
              penalty='l2',  
              l1_ratio=0.15)
```

Regularization parameters

Penalty values: 'l1', 'l2', 'elasticnet'

L1 optimiz.  
Function:

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 - \alpha \sum_{j=1}^k |\beta_j|$$


L2 optimiz.  
Function:

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2 - \alpha \sum_{j=1}^k \beta_j^2$$

## Mini Batch Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Combines of the best of both worlds ('Vanilla' Gradient Descent & Stochastic Gradient Descent): performs an update for every n training examples.




## Mini Batch Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Combines of the best of both worlds ('Vanilla' Gradient Descent & Stochastic Gradient Descent): performs an update for every n training examples.

$$w_1 = w_0 - \alpha \nabla 1/2 \sum_{i=1}^n \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$




# Mini Batch Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Combines of the best of both worlds ('Vanilla' Gradient Descent & Stochastic Gradient Descent): performs an update for every n training examples.

$$w_2 = w_1 - \alpha \nabla 1/2 \sum_{i=1}^n \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$




## Mini Batch Gradient Descent

$$J(\beta_0, \beta_1) = 1/2 \sum_{i=1}^m \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$

Combines of the best of both worlds ('Vanilla' Gradient Descent & Stochastic Gradient Descent): performs an update for every n training examples.

$$w_3 = w_2 - \alpha \nabla 1/2 \sum_{i=1}^n \left( (\beta_0 + \beta_1 x_{obs}^{(i)}) - y_{obs}^{(i)} \right)^2$$





## Mini Batch Gradient Descent

- Mini Batch implementation is typically used for neural nets
- Mini Batch sizes typically range from 50 to 256.
- There is a trade off between MB size and the learning rate ( $\alpha$ ): we can reduce learning rate for larger mini batch sizes (vice versa)
- We can tailor a learning rate schedule (i.e.): we can reduce the learning rate as go further along within our training epoch.
- We can implement with SGDClassifier (using 'partial fit') ex:  
`SGDClassifier(loss='log').partial_fit()`