

Errors in classification



		Condition (as determined by "Gold standard")			
Total population		Condition positive	Condition negative	Prevalence = $\frac{\Sigma \text{ Condition positive}}{\Sigma \text{ Total population}}$	
Test outcome	Test outcome positive	True positive	False positive (Type I error)	Positive predictive value (PPV, Precision) = $\frac{\Sigma \text{ True positive}}{\Sigma \text{ Test outcome positive}}$	False discovery rate (FDR) = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Test outcome positive}}$
	Test outcome negative	False negative (Type II error)	True negative	False omission rate (FOR) = $\frac{\Sigma \text{ False negative}}{\Sigma \text{ Test outcome negative}}$	Negative predictive value (NPV) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Test outcome negative}}$
Positive likelihood ratio (LR+) = TPR/FPR		True positive rate (TPR, Sensitivity, Recall) = $\frac{\Sigma \text{ True positive}}{\Sigma \text{ Condition positive}}$	False positive rate (FPR, Fall-out) = $\frac{\Sigma \text{ False positive}}{\Sigma \text{ Condition negative}}$	Accuracy (ACC) = $\frac{\Sigma \text{ True positive} + \Sigma \text{ True negative}}{\Sigma \text{ Total population}}$	
Negative likelihood ratio (LR-) = FNR/TNR		False negative rate (FNR) = $\frac{\Sigma \text{ False negative}}{\Sigma \text{ Condition positive}}$	True negative rate (TNR, Specificity, SPC) = $\frac{\Sigma \text{ True negative}}{\Sigma \text{ Condition negative}}$		
Diagnostic odds ratio (DOR) = LR+/LR-					

%54 Democrats, %46 republicans

Classify using their votes

%54 Democrats, %46 republicans

Classify using their votes

Model performance:

“How many times did I get it right?”

%54 Democrats, %46 republicans

Classify using their votes

Model performance:

“How many times did I get it right?”

Accuracy: % correct prediction of all predictions

%54 Democrats, %46 republicans

Classify using their votes

Model performance:

“How many times did I get it right?”

Accuracy: % correct prediction of all predictions

95% accuracy: Good job!

%1 have leukemia, %99 are healthy

Classify using health records and tests

%1 have leukemia, %99 are healthy

Classify using health records and tests

“How many times did I get it right?”

Accuracy: % correct prediction of all predictions

%1 have leukemia, %99 are healthy

Classify using health records and tests

“How many times did I get it right?”

Accuracy: % correct prediction of all predictions

Imagine the stupidest predictor:

“Always guess healthy”

%1 have leukemia, %99 are healthy

Classify using health records and tests

“How many times did I get it right?”

Accuracy: % correct prediction of all predictions

Imagine the stupidest predictor:

“Always guess healthy”

What will the accuracy be?

%1 have leukemia, %99 are healthy

Classify using health records and tests

“How many times did I get it right?”

Accuracy: % correct prediction of all predictions

Imagine the stupidest predictor:

“Always guess healthy”

What will the accuracy be?

It will be right 99% of the time!

You won't catch any sick people. Useless.

Confusion Matrix

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)
Spam (Actual)	27	6
Non-Spam (Actual)	10	57

Recall (Sensitivity) = $TP / (TP + FN)$

Precision = $TP / (TP + FP)$

Specificity = $TN / (TN + FP)$

Accuracy = $(TP + TN) / (TP + TN + FP + FN)$

Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)
Spam (Actual)	27	6
Non-Spam (Actual)	10	57

Recall (Sensitivity) = $TP / (TP + FN) = .82$

Precision = $TP / (TP + FP) = .73$

Specificity = $TN / (TN + FP) = .85$

Accuracy = $(TP + TN) / (TP + TN + FP + FN) = .84$

Precision and recall

Precision and recall

Precision: Out of all cases I predicted as positive,
how many times was I right?

Precision and recall

Precision: Out of all cases I predicted as positive,
how many times was I right?

(% times I was right when I told somebody they
had leukemia)

Precision and recall

Precision: Out of all cases I predicted as positive,
how many times was I right?

(% times I was right when I told somebody they
had leukemia)

Recall: Out of all the (few) positive cases,
how many did I find

Precision and recall

Precision: Out of all cases I predicted as positive,
how many times was I right?

(% times I was right when I told somebody they
had leukemia)

Recall: Out of all the (few) positive cases,
how many did I find

(% of actual leukemia patients I could catch with
my classifier)

Confusion Matrix

	Spam (Predicted)	Non-Spam (Predicted)
Spam (Actual)	0	10
Non-Spam (Actual)	0	990

Recall (Sensitivity) = $TP / (TP + FN) = 0/10 = 0$

Precision = $TP / (TP + FP) = 0/0 \rightarrow \text{undefined!}$

Specificity = $TN / (TN + FP) = 100\%$

Accuracy = $(TP + TN) / (TP + TN + FP + FN) = 99\%$

Confusion Matrix

	p' (Predicted)	n' (Predicted)
p (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

Confusion Matrix

	p' (Predicted)	n' (Predicted)
P (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Focusing on a
single class
(positive: the
one with small
prevalence) in
skewed cases

$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

Confusion Matrix

	p' (Predicted)	n' (Predicted)
P (Actual)	True Positive	False Negative
n (Actual)	False Positive	True Negative

Focusing on a
single class
(positive: the
one with small
prevalence) in
skewed cases

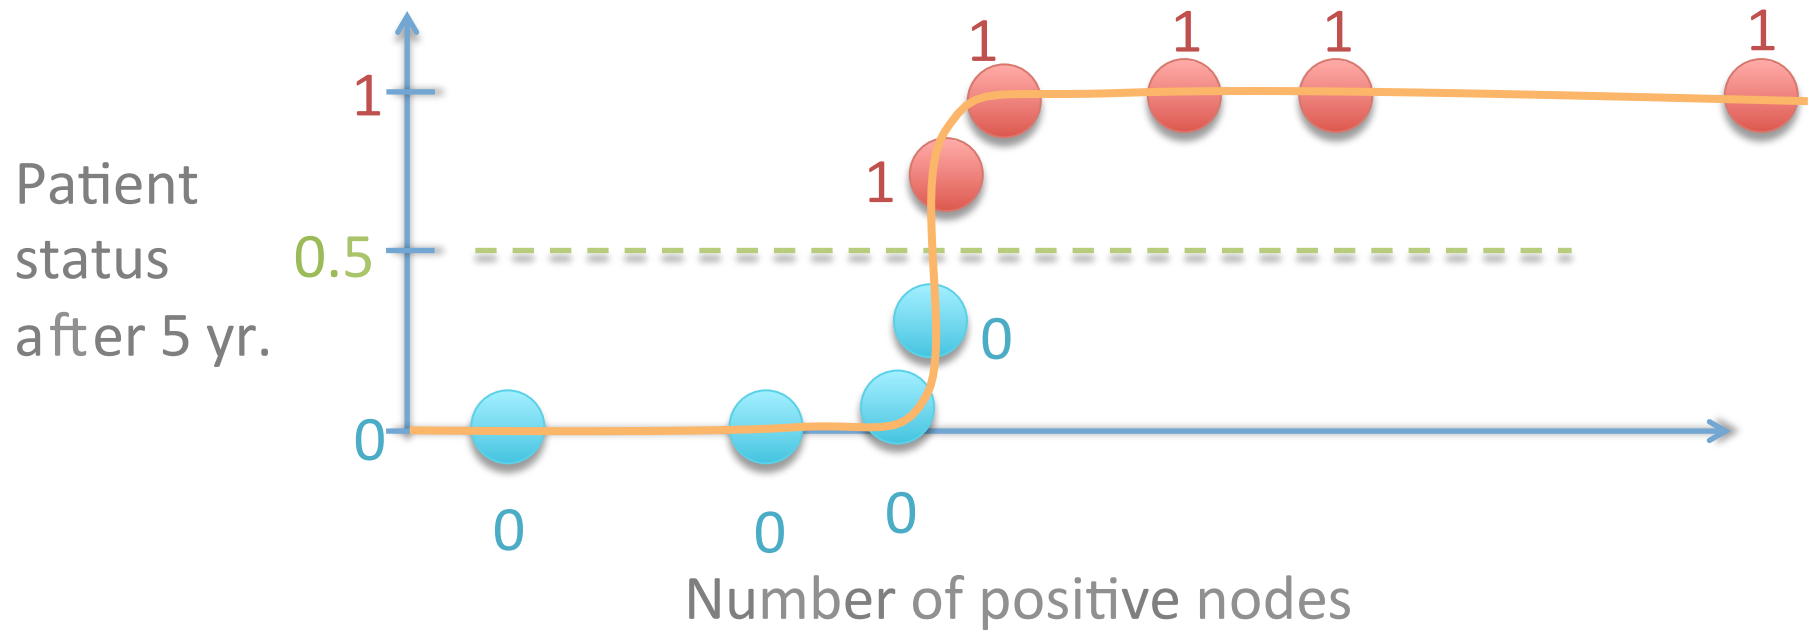
$$\text{Precision} = \text{TP} / (\text{TP} + \text{FP})$$

$$\text{Recall} = \text{TP} / (\text{TP} + \text{FN})$$

F1 = Their harmonic mean


$$F1 = 2 * \frac{\text{precision} * \text{recall}}{\text{precision} + \text{recall}}$$

Logistic regression

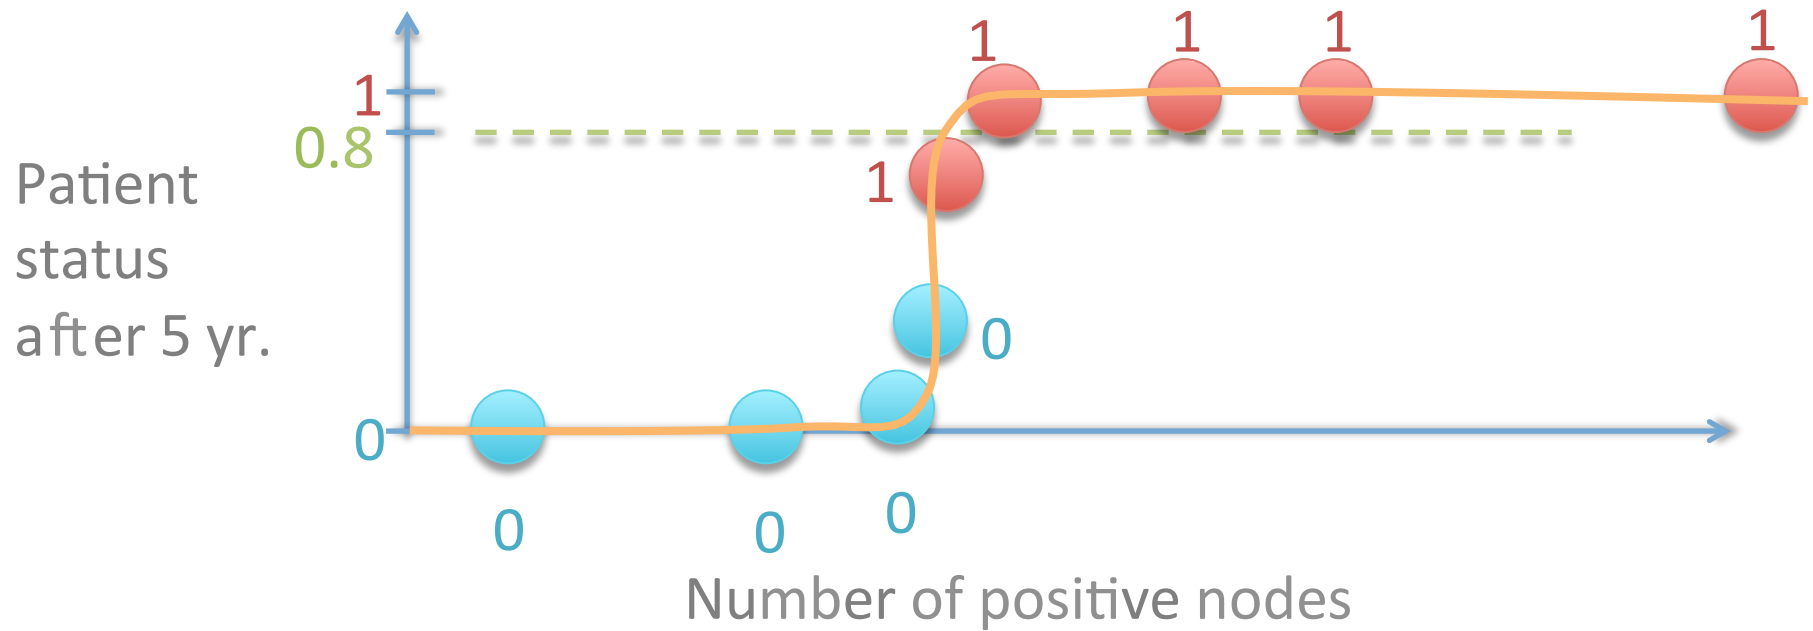


$$y_{\beta}(x) = \frac{1}{1 + e^{-(\beta_0 + \beta_1 x + \varepsilon)}}$$

Logistic regression



Logistic regression




Higher threshold: More sure about positives

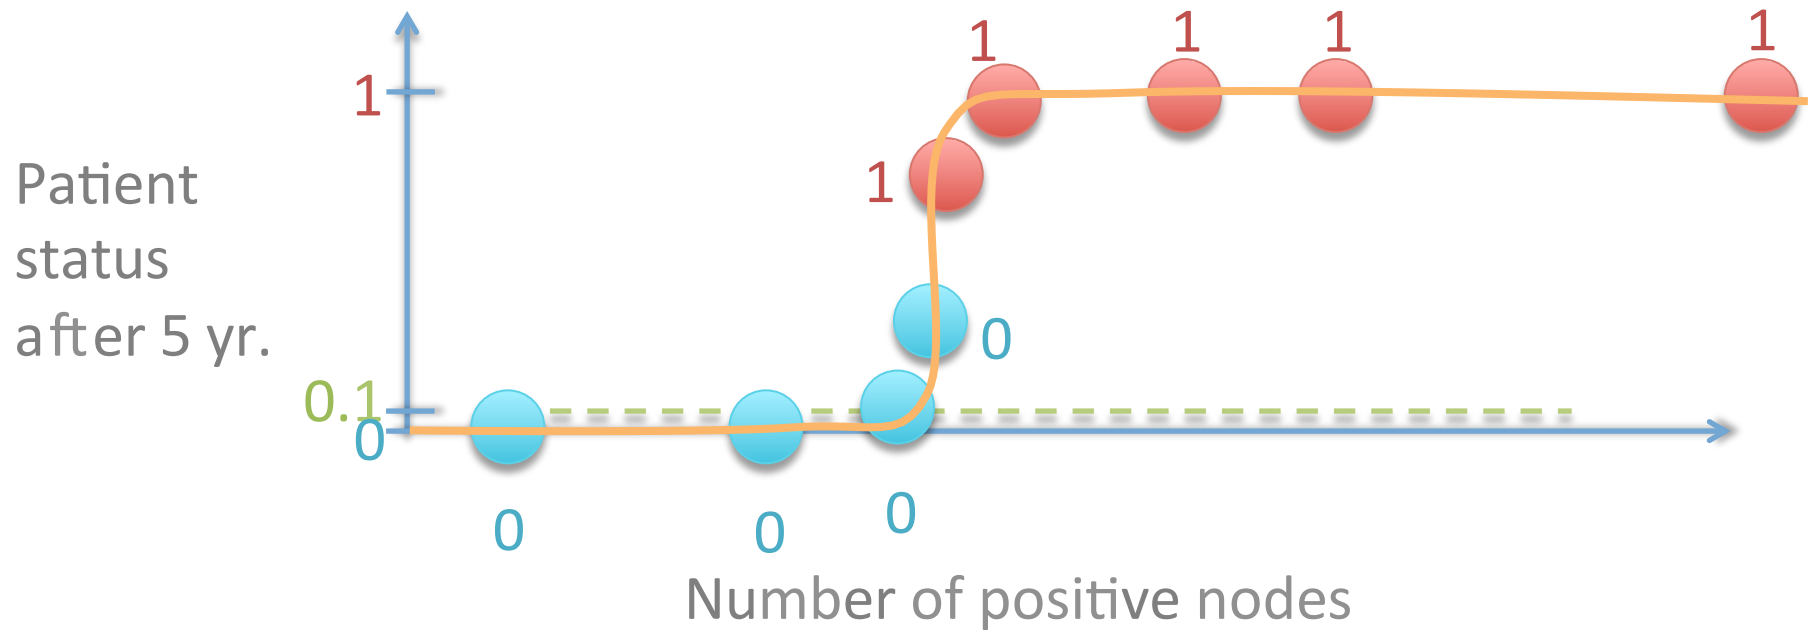
lower recall, higher precision

lower True Positive Rate, lower False Positive Rate

Logistic regression



Logistic regression



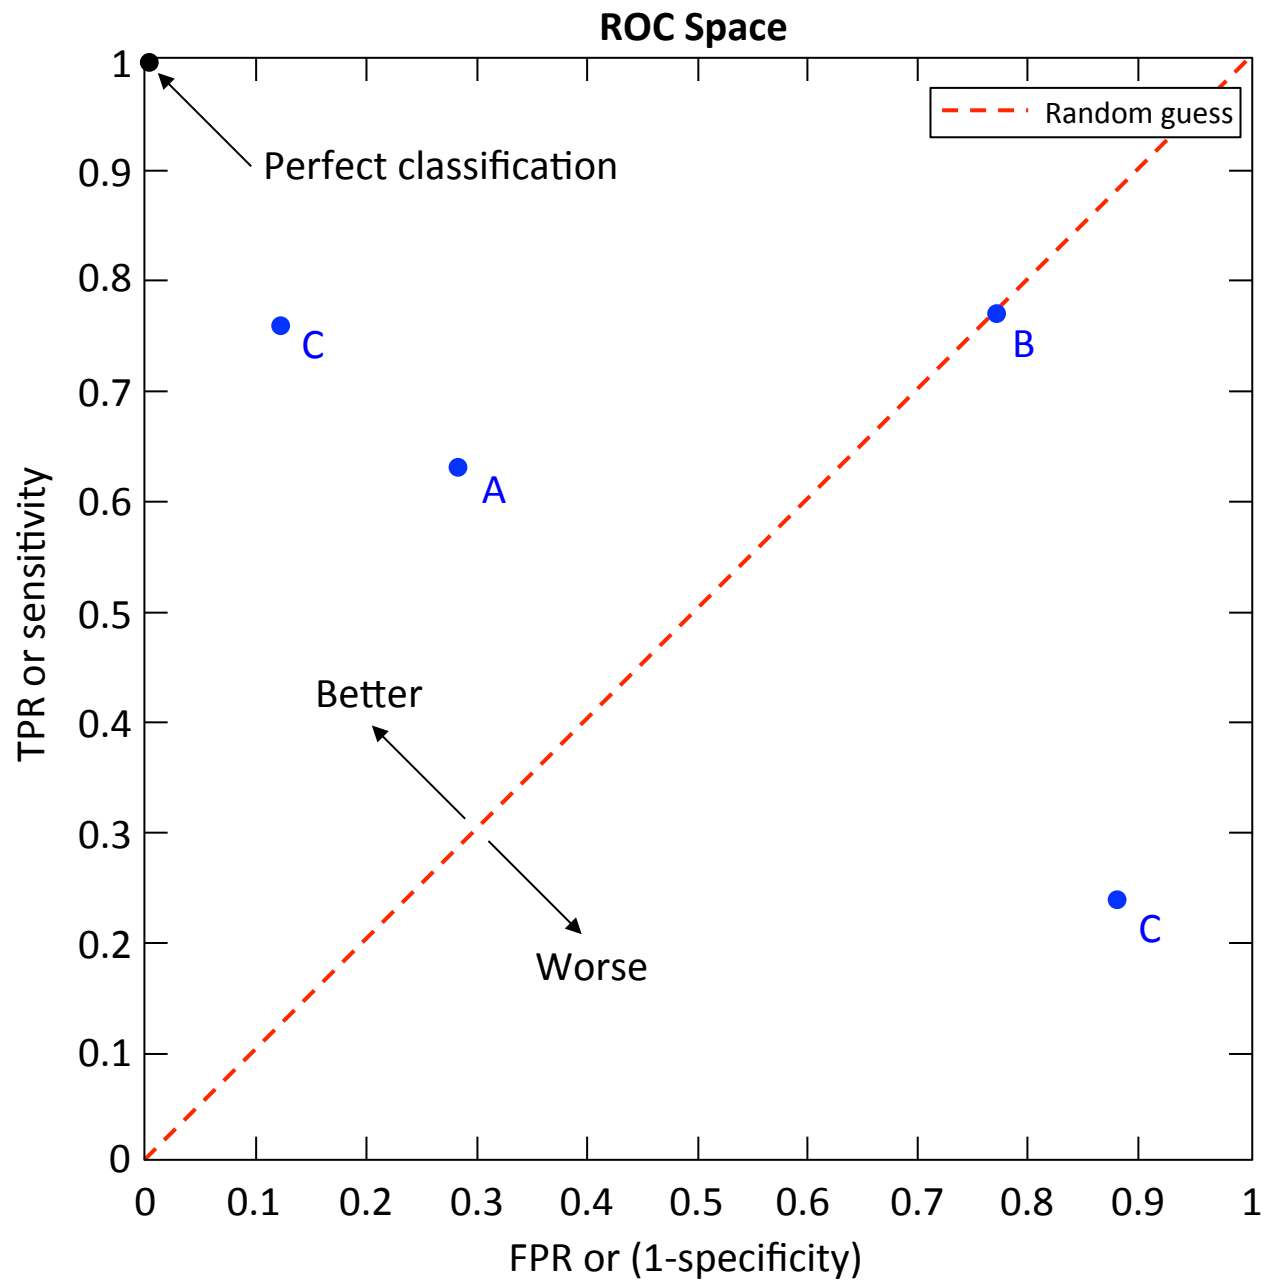
Lower threshold: Better at catching positives

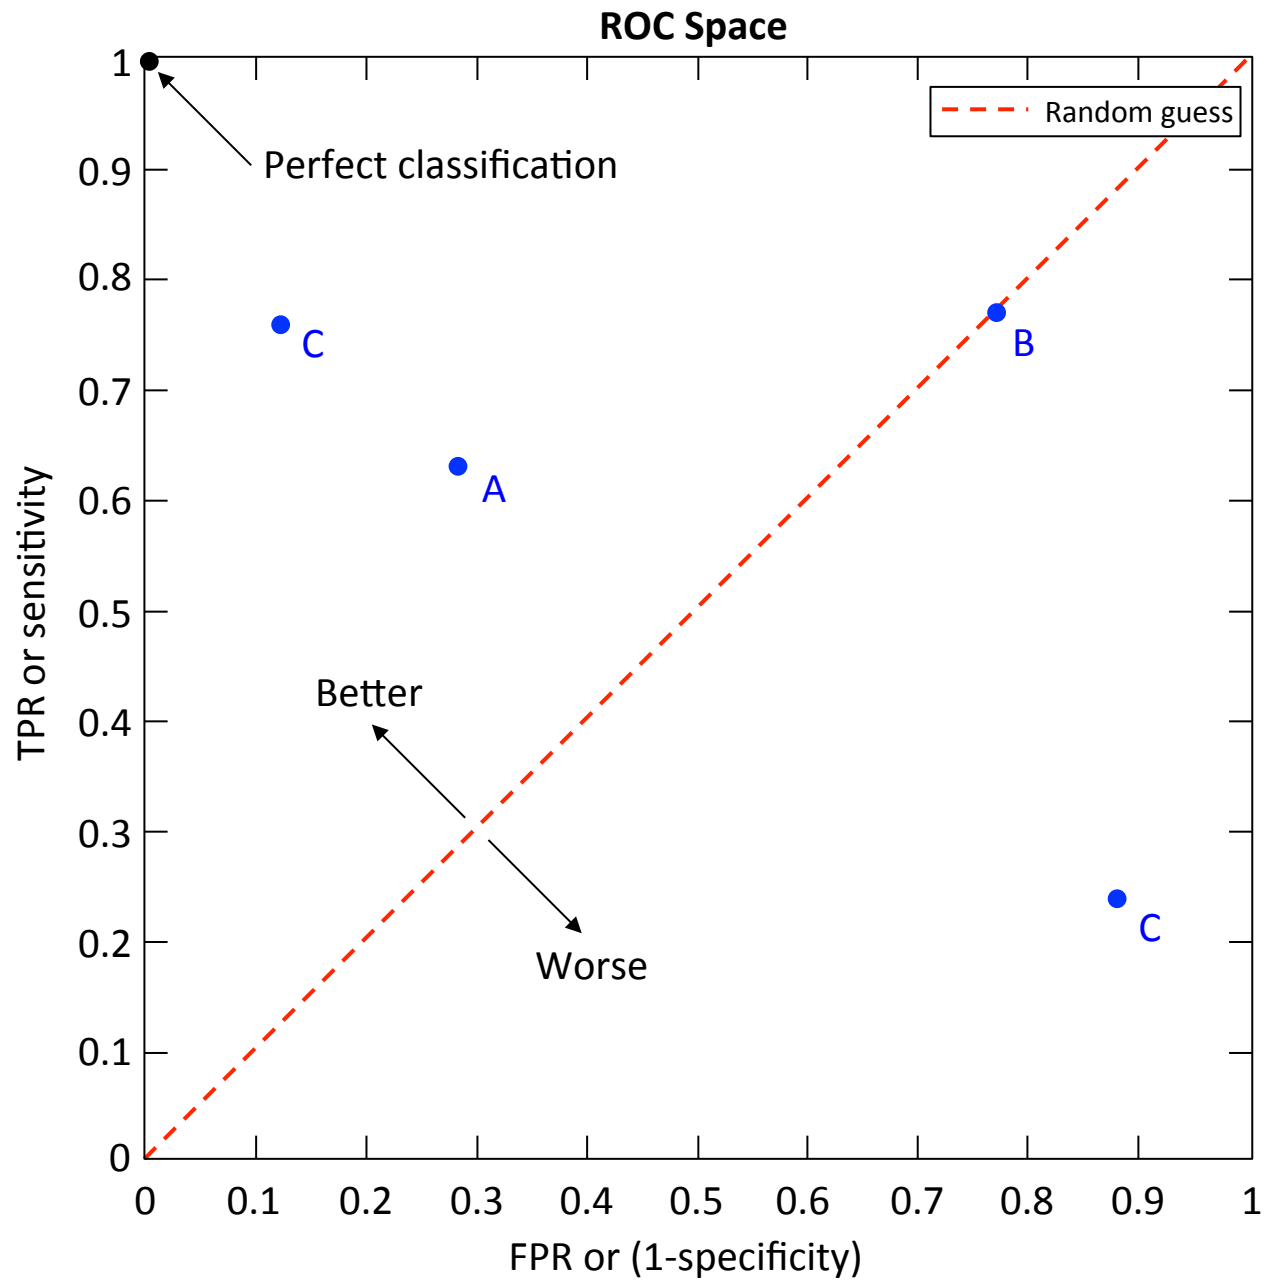
higher recall, less precision

higher True Positive Rate, higher False Positive Rate

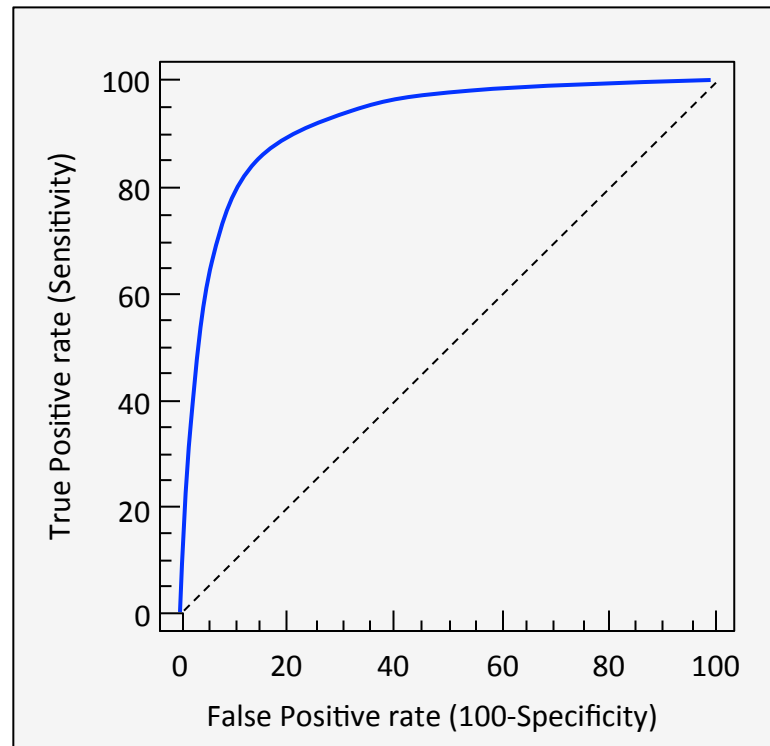
Each threshold is a different model

Plot their True Positive Rate & False Positive Rate





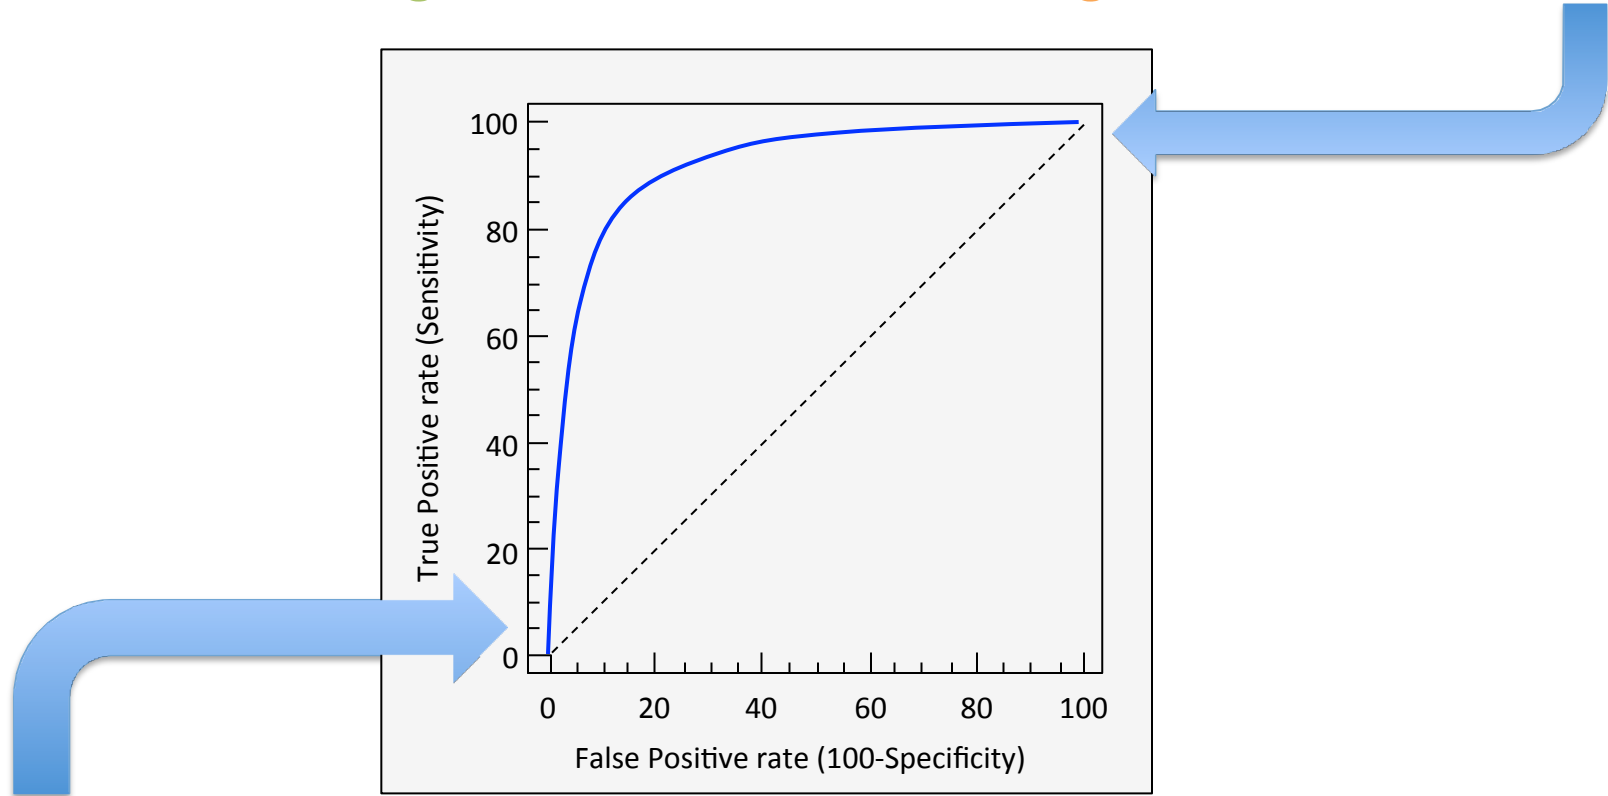
Receiver Operating Characteristic



Lower threshold: Better at catching positives

higher recall, less precision

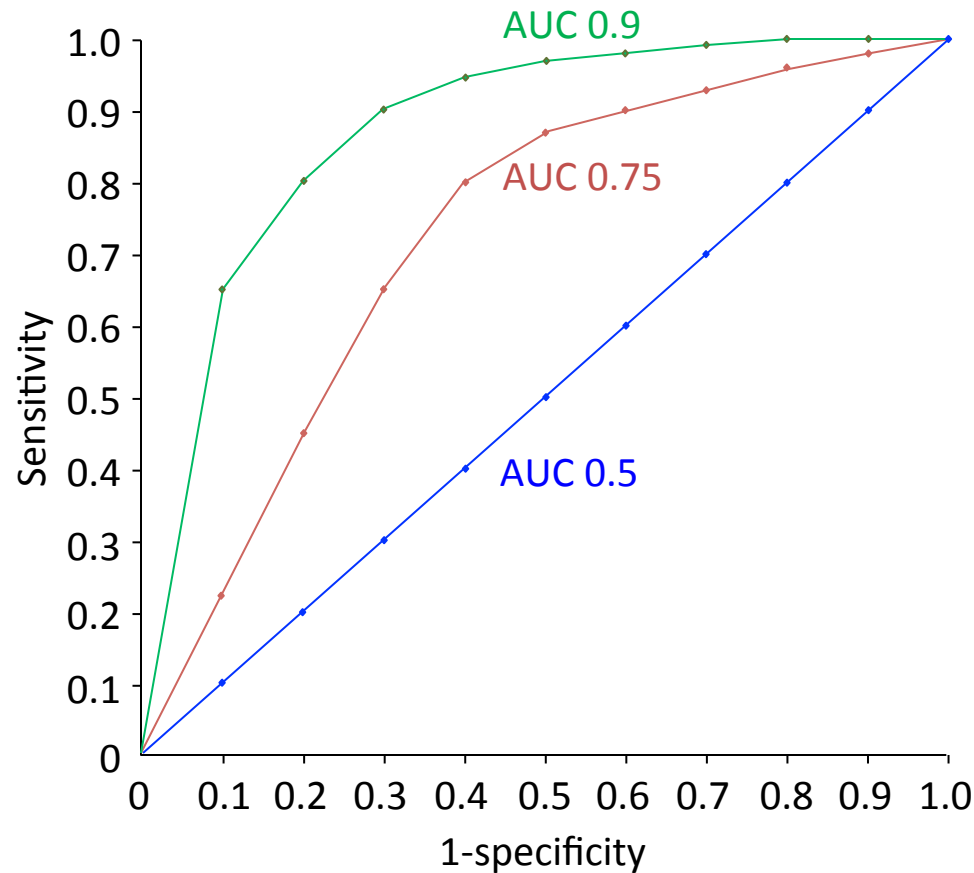
higher True Positive Rate, higher False Positive Rate



Higher threshold: More sure about positives

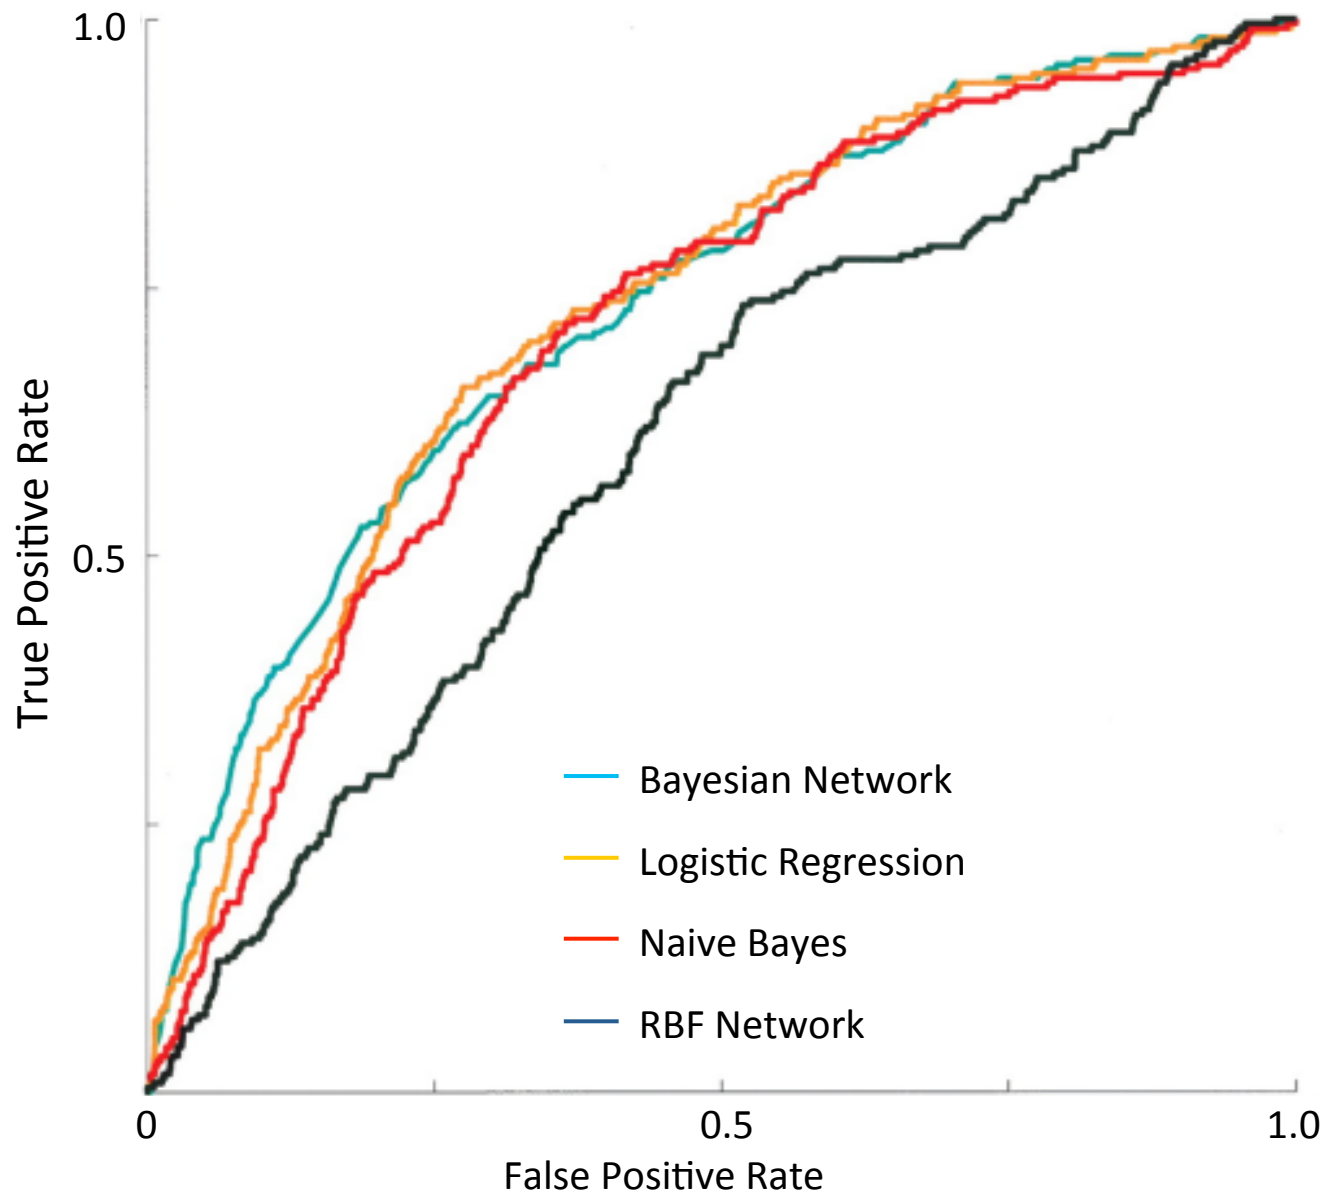
lower recall, higher precision

lower True Positive Rate, lower False Positive Rate



Area under curve (AUC)

An evaluation of a classification algorithm
(including all possible thresholds)



from sklearn.metrics import

Classification metrics

See the [Classification metrics](#) section of the user guide for further details.

<code>metrics.accuracy_score(y_true, y_pred[, ...])</code>	Accuracy classification score.
<code>metrics.auc(x, y[, reorder])</code>	Compute Area Under the Curve (AUC) using the trapezoidal rule
<code>metrics.average_precision_score(y_true, y_score)</code>	Compute average precision (AP) from prediction scores
<code>metrics.classification_report(y_true, y_pred)</code>	Build a text report showing the main classification metrics
<code>metrics.confusion_matrix(y_true, y_pred[, ...])</code>	Compute confusion matrix to evaluate the accuracy of a classification
<code>metrics.f1_score(y_true, y_pred[, labels, ...])</code>	Compute the F1 score, also known as balanced F-score or F-measure
<code>metrics.fbeta_score(y_true, y_pred, beta[, ...])</code>	Compute the F-beta score
<code>metrics.hamming_loss(y_true, y_pred[, classes])</code>	Compute the average Hamming loss.
<code>metrics.hinge_loss(y_true, pred_decision[, ...])</code>	Average hinge loss (non-regularized)
<code>metrics.jaccard_similarity_score(y_true, y_pred)</code>	Jaccard similarity coefficient score
<code>metrics.log_loss(y_true, y_pred[, eps, ...])</code>	Log loss, aka logistic loss or cross-entropy loss.
<code>metrics.matthews_corrcoef(y_true, y_pred)</code>	Compute the Matthews correlation coefficient (MCC) for binary classes
<code>metrics.precision_recall_curve(y_true, ...)</code>	Compute precision-recall pairs for different probability thresholds
<code>metrics.precision_recall_fscore_support(...)</code>	Compute precision, recall, F-measure and support for each class
<code>metrics.precision_score(y_true, y_pred[, ...])</code>	Compute the precision
<code>metrics.recall_score(y_true, y_pred[, ...])</code>	Compute the recall
<code>metrics.roc_auc_score(y_true, y_score[, ...])</code>	Compute Area Under the Curve (AUC) from prediction scores
<code>metrics.roc_curve(y_true, y_score[, ...])</code>	Compute Receiver operating characteristic (ROC)
<code>metrics.zero_one_loss(y_true, y_pred[, ...])</code>	Zero-one classification loss.

Always remember,

Fit the model to a **training set**,

Calculate performance

(**accuracy**, **precision**, **recall**, **f1**, **AUC**, etc.)

on a **test set**

or (better) on a k-fold **cross validation** scheme