


# Support Vector Machines

Summer 2018



# Machine learning explained by XKCD



<https://xkcd.com/1838/>

# What is a support vector machine?


Supervised, linear classification

You, typeface connoisseur: Is that Comic Sans?!

Me, science genius: I read some paper on the internet once that people remember things written in ugly typefaces better...

# SVMs look like a line in 2D space

X2 = fans who know the South Side is the best side



Support vector machine hyperplane

X1 = fans who love to see their team lose

# Revisiting our machine learning schema

- SVMs are a type of **supervised learning model**
- Can be used both for **classification and regression** (covering only classification today)
- They are a **binary classifier** (can be extended to multiple classes, but gets complicated)
- It is a **linear classifier**, meaning it uses a linear combination of the inputs to make its classification prediction
- Other ways of doing supervised classification include: logistic regression, k-nearest neighbors, decision trees, neural networks

# Reasons to care about support vector machines

1. Someone will ask you about it in an interview
2. It's easy to interpret when you need a binary classifier
3. It's memory efficient
4. It's robust to a whole bunch of issues, including sparse data and high dimensions

# In the beginning, there was the perceptron




**Perceptron**

# In the beginning, there was the perceptron

## Perceptron

- Invented in late 1950s
- Separates observations into classes using **a hyperplane as the decision boundary**
- Requires that the classes can be separated by a line (**linear separability**)
- Rigid and sensitive to outliers




$$f(x) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \\ 0 & \text{otherwise} \end{cases}$$



# In the beginning, there was the perceptron

## Perceptron


- Invented in late 1950s
- Separates observations into classes using **a hyperplane as the decision boundary**
- Requires that the classes can be separated by a line (**linear separability**)
- Rigid and sensitive to outliers



$$f(x) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \\ 0 & \text{otherwise} \end{cases}$$

## Maximal margin classifier


- Makes statistically concrete the **intuition that some lines are better** than others
- Takes perceptron and adds **optimal stability** (widest margin)
- Fails completely if data not linearly separable



# In the beginning, there was the perceptron

## Perceptron


- Invented in late 1950s
- Separates observations into classes using **a hyperplane as the decision boundary**
- Requires that the classes can be separated by a line (**linear separability**)
- Rigid and sensitive to outliers



$$f(x) = \begin{cases} 1 & \text{if } \mathbf{w} \cdot \mathbf{x} + b > 0 \\ 0 & \text{otherwise} \end{cases}$$


## Maximal margin classifier

- Makes statistically concrete the **intuition that some lines are better** than others
- Takes perceptron and adds **optimal stability** (widest margin)
- Fails completely if data not linearly separable




## Support vector machine

- Developed early 1990s
- SVM adds two generalizations:
  1. **Soft margin** for outliers
  2. **Kernel trick** for data that's not **linearly separable**: gives us complex feature space with minimal computational complexity




SVMs can be further extended to multi-class data.

# What's the use of SVM? Splitting the feature space to enable prediction




- Why are we interested in finding the best line (hyperplane) that separates our classes?
- In order to predict new instances or observations, we need to separate our feature space
- A better split gives us **higher confidence** that our predictions of new data are correct

# Defining the decision boundary: the support vectors of SVM



- Each observation is represented by a vector of values
- Why do we define the margin by the distance to the closest points (and not, e.g., by the average distance)?
- SVM is insensitive to points far away from the decision boundary
- The separating hyperplane is defined by the **support vectors** (highlighted in orange)
- If there are no degeneracies (if the data is linearly separable), then we **need  $d+1$  support vectors** (where  $d$  is the # of dimensions)

# Creating the largest margin between classes (hard margin)



- Why do we want the largest margin?
  - **Intuition:** the larger the margin, the more confidence we have that we have the correct classification of our test data
  - **Statistics:** the VC dimension, or how complex a model is and how likely it is to capture too many wrong points
  - E.g. a curved line may fit the test data better but is more likely to make mistakes in prediction

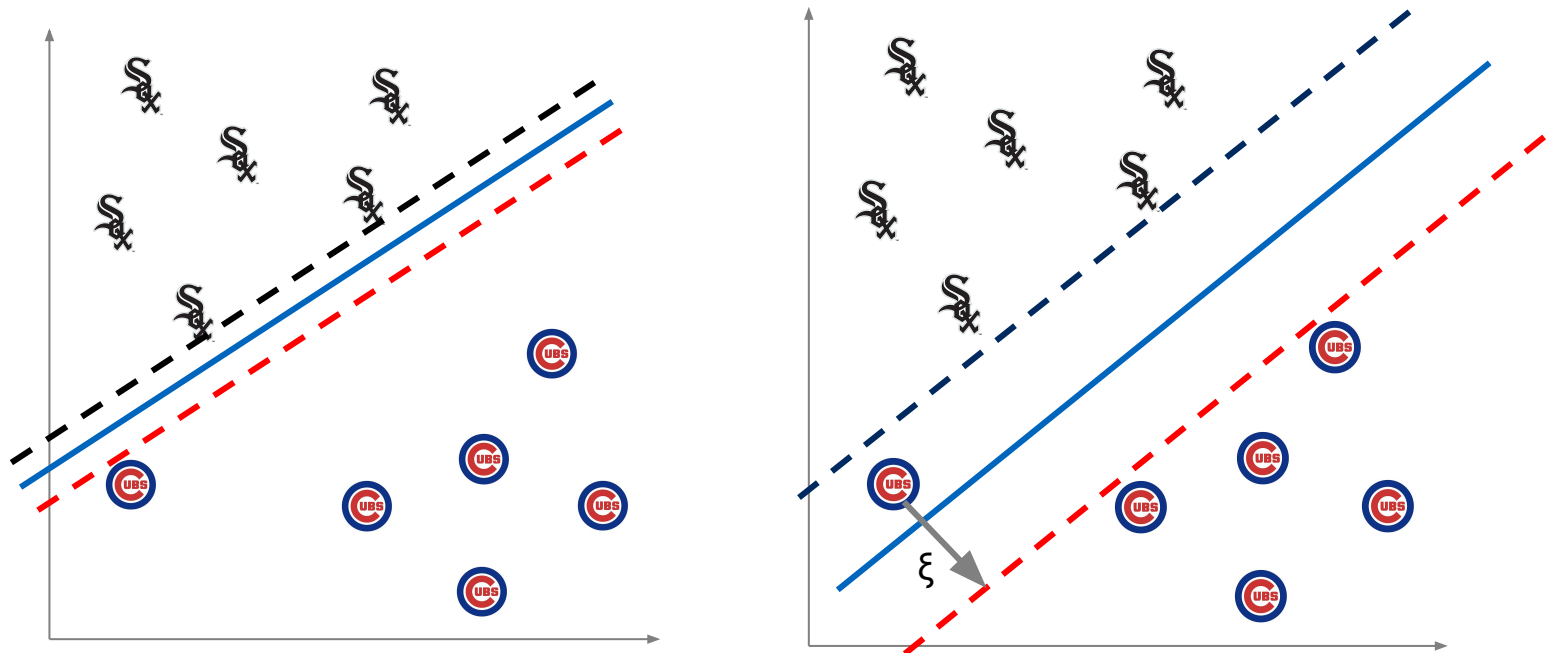
## Solving for the largest margin

$$\begin{array}{ll} \max_{w, \gamma} & \gamma \\ \text{s.t.} & \forall i, y_i (w \cdot x_i + b) \geq \gamma \end{array} \quad \longrightarrow \quad \begin{array}{ll} \min_w & \frac{1}{2} \|w\|^2 \\ \text{s.t.} & \forall i, y_i (w \cdot x_i + b) \geq 1 \end{array}$$

- Solving for the margin is a **constrained optimization** problem
- If we maximize  $\gamma$ , we can do so by increasing  $w$  (weights) as much as we want
- Instead, rewrite to maximize the unit  $w$
- SVM can be used for **online learning**, where the model parameters are slightly modified with the introduction of a new data point and where data comes in as a stream


## From maximal margin classifier to SVM: $\xi$ and kernel trick

# SVM: Soft margin classification




- A **hard margin** (1) works only if the data is linearly separable, and (2) is sensitive to outliers as it will work to completely capture the data on the correct side of the line
- We create a **soft margin** by adding in **slack to our cost function**
- We manually set **C**, a tuning parameter, to tell the algorithm how much slack it has to misclassify observations
  - Large C = misclassify more

# SVM: Using the kernel trick when the data isn't linearly separable




- When the classes aren't linearly separable, we can't use a linear separator
- In the regression case (e.g. OLS), we dealt with this by adding to the dimension of our feature space by transforming the regressors (e.g.  $x^2$ )
- The kernel trick is a similar approach
- **Mathematical aside:** turns out, to solve our SVM, we need only the dot product of the observations (instead of the observations themselves)
- **In the kernel trick, replace every instance of the dot product with a new function** (e.g. polynomial, like  $X^2$ ) that's equivalent to the dot product






# Visualizing the kernel trick

Best case scenario is when our data is linearly separable: it's easy to draw our separating line in this case:




# Visualizing the kernel trick

Best case scenario is when our data is linearly separable: it's easy to draw our separating line in this case:




Are we out of luck if our data isn't perfectly separable?




# Visualizing the kernel trick


Best case scenario is when our data is linearly separable: it's easy to draw our separating line in this case:




Are we out of luck if our data isn't perfectly separable?



We can still use a linear separator, if we shift our data into a higher-dimensional space!




# Visualizing the kernel trick




# SVM: Using the kernel trick when the data isn't linearly separable

Polynomial kernel (=3)



Radial kernel





- **Mathematical aside:** turns out, to solve our SVM, we need only the inner product of the observations (instead of the observations themselves)
- In the kernel trick, replace every instance of the inner product with a new function (e.g. polynomial, like  $X^2$ )
- But, **why use kernels instead of just enlarging the feature space directly?**
  - By using the kernel, we need to compute only the dot product of all observations, not the actual transformations
  - This **saves computation**
  - Some expansions are infinite so we couldn't solve them without the kernel anyway

## Tuning SVM hyperparameters: $C$ and $\gamma$



# Tuning the hyperparameters of an SVM: soft margin through “C”

Largest value of C (low variance)



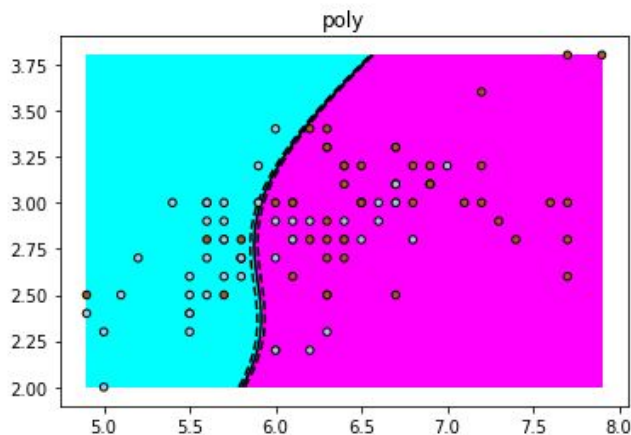
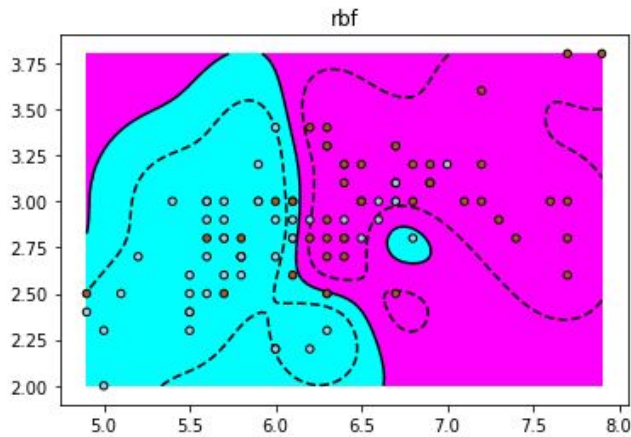
- The level of **C** tells the algorithm **how much slack it has to misclassify** some observations
- Finding the **optimal value** of C requires manual tuning and in practice is **often done via grid search and cross-validation**

Smallest value of C (low bias)





# Tuning the hyperparameters of an SVM: complexity through $\gamma$

$\gamma=10$



Increase  
complexity

$\gamma=20$  (more complex)




- The hyperparameter gamma is **used only with non-linear kernels** (e.g. polynomial, radial basis function (RBF))
- **Gamma** tells the model the kernel coefficient, which **affects how each observation influences the support vectors**







## Extending to multiclass

# Extension to multiclass estimating: one vs. rest

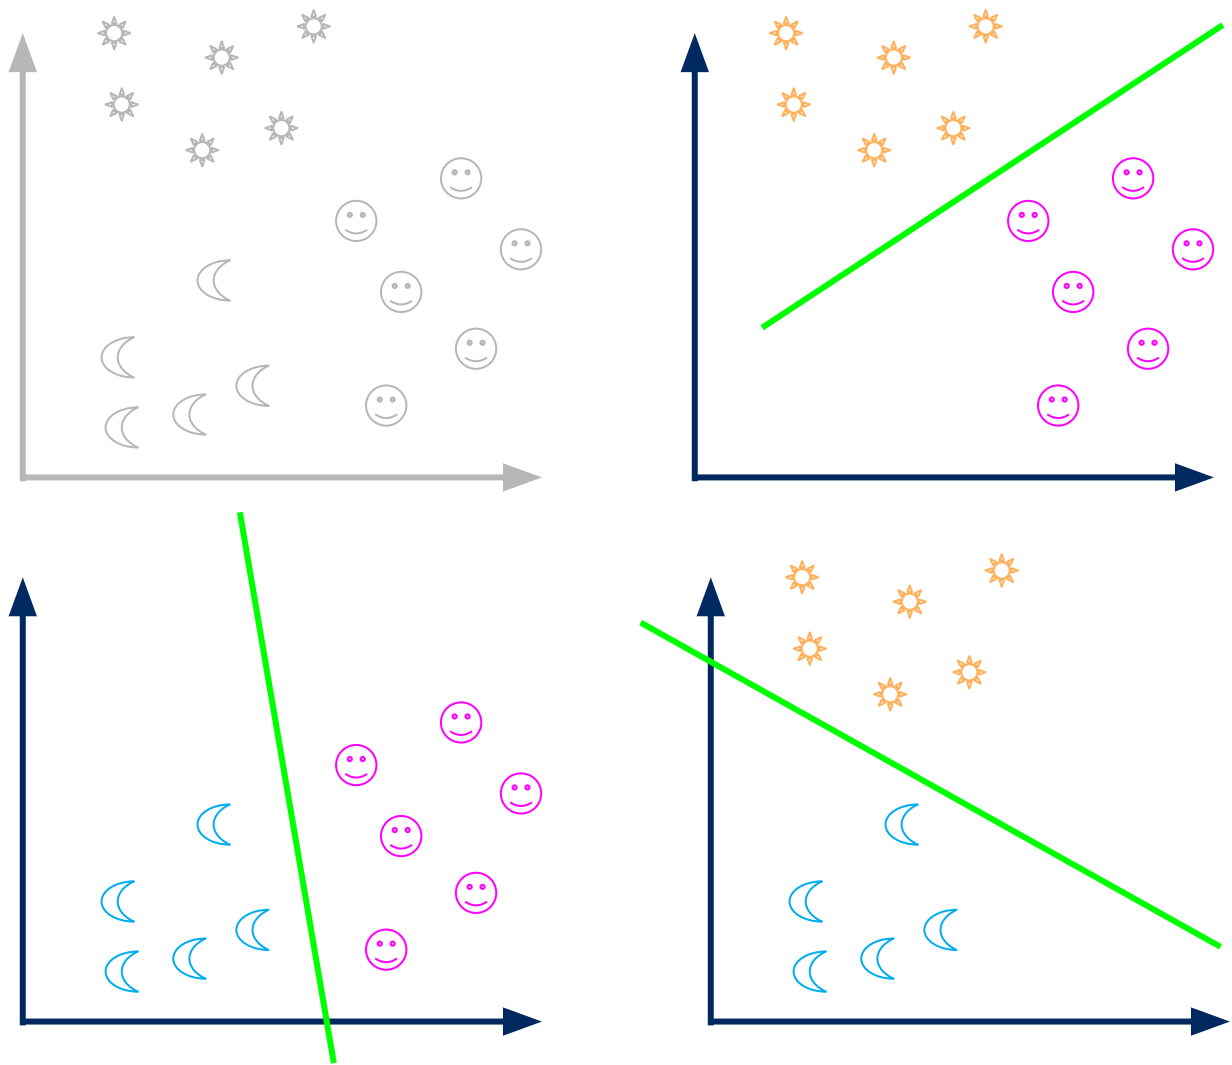


# Extension to multiclass estimating: one vs. rest



Three models are fit, and prediction is done using the model that would give the widest margin.


# Extension to multiclass estimating: one vs. one



Pairwise models are fit, and prediction is done using majority voting.


# Expanding to multiple classes

## OVR: One vs Rest



**Pros: Fewer classifications**  
**Cons: Classes may be imbalanced**

## OVO: One vs One



**Pros: Less sensitive to imbalance**  
**Cons: More classifications**

# Scaling data before using SVM

- Inputs should be scaled prior to running SVM
- Most common scaling: subtract the mean and divide by the SD
  - Remember to scale the testing dataset with the mean and SD of the training dataset
- Why scale?
  - Avoid data with large ranges dominating columns with smaller ranges
  - Avoid numerical issues from kernel computation

# When should we use an SVM?

## Pros

- Easy to interpret (output is a class)
- Can be used with high dimensional data, even if the # dimensions > # observations
- Can be used with sparse data
- Robust to outliers that are far away from the decision hyperplane (which is affected only by the support vectors)
- Defined by low number of support vectors so memory efficient

## Cons

- Does not compute probabilities
- To compute probabilities requires cross-validation, which is computationally expensive
- Perform poorly with unbalanced classes
- Performs poorly with overlapping classes
- Hard to generalize beyond two classes (though possible)

# Implementing SVM in Python with scikit-learn

## Fitting a model

```
class sklearn.svm. svc (C=1.0, kernel='rbf', degree=3, gamma='auto', coef0=0.0, shrinking=True, probability=False,  
tol=0.001, cache_size=200, class_weight=None, verbose=False, max_iter=-1, decision_function_shape='ovr',  
random_state=None)
```



# Implementing SVM in Python with scikit-learn

## Fitting a model and predicting a new observation


```
In [2]: # create 40 separable points in two classes
X, y = make_blobs(n_samples=40, centers=2, random_state=6)

# fit the model in two steps

# 1) Set the hyperparameters
model = svm.SVC(kernel='linear', C=1000)

#2) Fit the model
model.fit(X, y)
```

```
Out[2]: SVC(C=1000, cache_size=200, class_weight=None, coef0=0.0,
  decision_function_shape='ovr', degree=3, gamma='auto', kernel='linear',
  max_iter=-1, probability=False, random_state=None, shrinking=True,
  tol=0.001, verbose=False)
```



To predict a new data point, we test (6,11) and (5,0):

```
model.predict(np.array([6,11]).reshape(1,-1))
array([0])
```

```
model.predict(np.array([5,0]).reshape(1,-1))
array([1])
```