

DEEP LEARNING ARCHITECTURES

A QUICKSTART GUIDE
TO THE NUTS & BOLTS
OF DEEP LEARNING

► By: METIS



DEEP LEARNING OVERVIEW

- ▶ Deep Learning = Fancy ANNs (Artificial Neural Networks)
 - ▶ **Deep** means many hidden layers... think “patterns of patterns”
 - ▶ **Fancy** means complex network structures that capture more information from data
- ▶ Deep Learning neural networks are universal function approximators. But there is no free lunch! We pay for this power with having a **massive** number of parameters and hyperparameters to tune.
 - ▶ https://en.wikipedia.org/wiki/Universal_approximation_theorem



DEEP LEARNING APPLICATIONS

- ▶ Machine Translation
- ▶ Named Entity Recognition
- ▶ Image Processing (see CNNs lecture)
- ▶ Image + Video Captioning
- ▶ Chatbots
- ▶ Text Generation
- ▶ Question Answering



BIAS / VARIANCE TRADEOFF

- ▶ Bias = error from erroneous assumptions in the learning algorithm
- ▶ Variance = error from sensitivity to small fluctuations in training data
- ▶ Generalization = test performance vs. training performance
- ▶ As usual, larger training sets, dimensionality reduction, and feature selection can decrease variance by simplifying models.
- ▶ In neural networks, variance increases and bias decreases as the number of hidden units increases.
- ▶ Regularization can be applied via Dropout (we'll see this soon).



FEED FORWARD / BACK PROP

The concepts in any Deep Learning model are all the same:

- ▶ Nodes feedforward to their successors and backpropagate to update weights to their predecessors
- ▶ Weights matrices represent transitions between nodes
- ▶ Depth can actually save us a lot of parameters
 - ▶ For dense layers, number of parameters is $O(n^2)$ in nodes per layer, but $O(n)$ in number of layers.



SIMPLE CONCEPTS, ENDLESS OPTIONS

- ▶ Take the input
- ▶ Run it thru the network
- ▶ Compute the Error/Cost
- ▶ Backpropagate to update weights via gradient of cost function
- ▶ **Everything is Gradient Descent.**
- ▶ That being said, Deep Learning will drown you in options!



CHOOSE THE ARCHITECTURE

Layers and Activation functions determine what kinds of patterns the network will detect in your data.

- ▶ Number of Layers
- ▶ Types of Layers
- ▶ Number of Neurons per Layer
- ▶ Activation Function for each Layer



CHOOSE AN OPTIMIZER

Optimizer algorithms are just variations of Gradient Descent

- ▶ SGD
- ▶ SGD w/ momentum
- ▶ Adam
- ▶ Adagrad
- ▶ RMSprop



CHOOSE A LOSS FUNCTION

- ▶ Mean Squared Error
- ▶ Hinge
- ▶ Binary Cross Entropy (Log loss)
- ▶ Categorical Cross Entropy
- ▶ Many more



GENERAL WORKFLOW

- ▶ Find an architecture that meets or exceeds your performance goals on the training data (potentially overfitting)
- ▶ Introduce regularization (dropout) to force generalization
- ▶ If training is too slow:
 - ▶ use GPUs to speed up calculations
 - ▶ try transfer learning from pretrained model or from autoencoder






Deep Learning Layers

WE'RE JUST PLAYING WITH LEGOS FROM HERE ON OUT


DENSE LAYERS

- ▶ Simplest type of neural network layer to understand.
- ▶ Also known as “fully connected” layers.
- ▶ Every node in the output layer is a linear combination of each node in the input layer, passed through an activation function.
- ▶ Basically just a layer full of logistic regressions with different weights.




DROPOUT LAYERS

- ▶ Dropout is Deep Learning's take on Regularization.
- ▶ Weights are randomly ignored during training so that the remaining nodes are forced to learn useful information.



Standard Neural Net



After applying dropout

source: [Deep Learning for Computer Vision by Rajalingappaa Shanmugamani](#)




RECURRENT NEURAL NETWORKS

- ▶ So far, we've stuck to simple ANNs (fully connected)
- ▶ Recurrent Neural Networks (RNNs) change the game
- ▶ Connections between units form a **directed cycle**
- ▶ Allows network to retain internal state from past units → **memory**
- ▶ Allows dynamic temporal behavior
- ▶ Can use memory to process arbitrary input sequences!
- ▶ Terrific success in various NLP tasks



RECURRENT LAYERS


- ▶ How do they do it?
- ▶ Hidden Units at a time step t are dependent on:
 - ▶ The previous hidden unit
 - ▶ The input at time step t




RECURRENT LAYERS

- ▶ Drawback of ANNs: Fixed # of inputs and outputs
- ▶ RNNs can map arbitrary length sequences of each.
- ▶ Don't stress! Just a different architecture with some nice features. Same solving concepts apply. (Keras takes care of this for you anyway!)


one to one




one to many




many to one



many to many

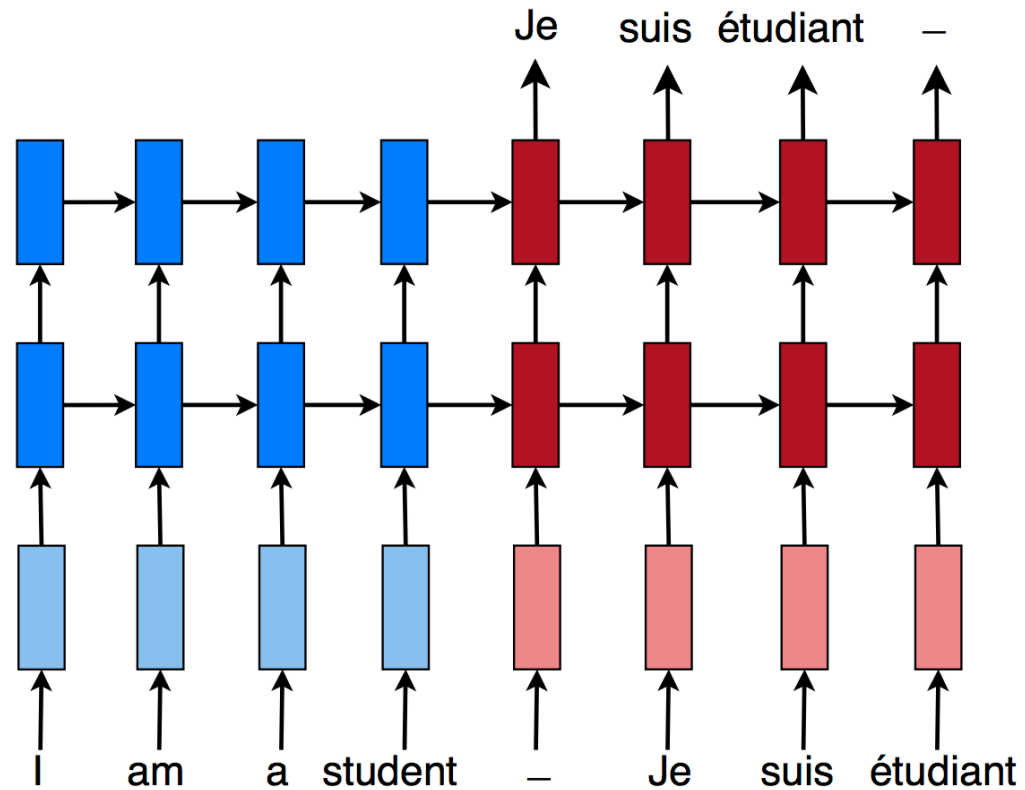


many to many



RECURRENT LAYERS

- Sequence-to-sequence mapping

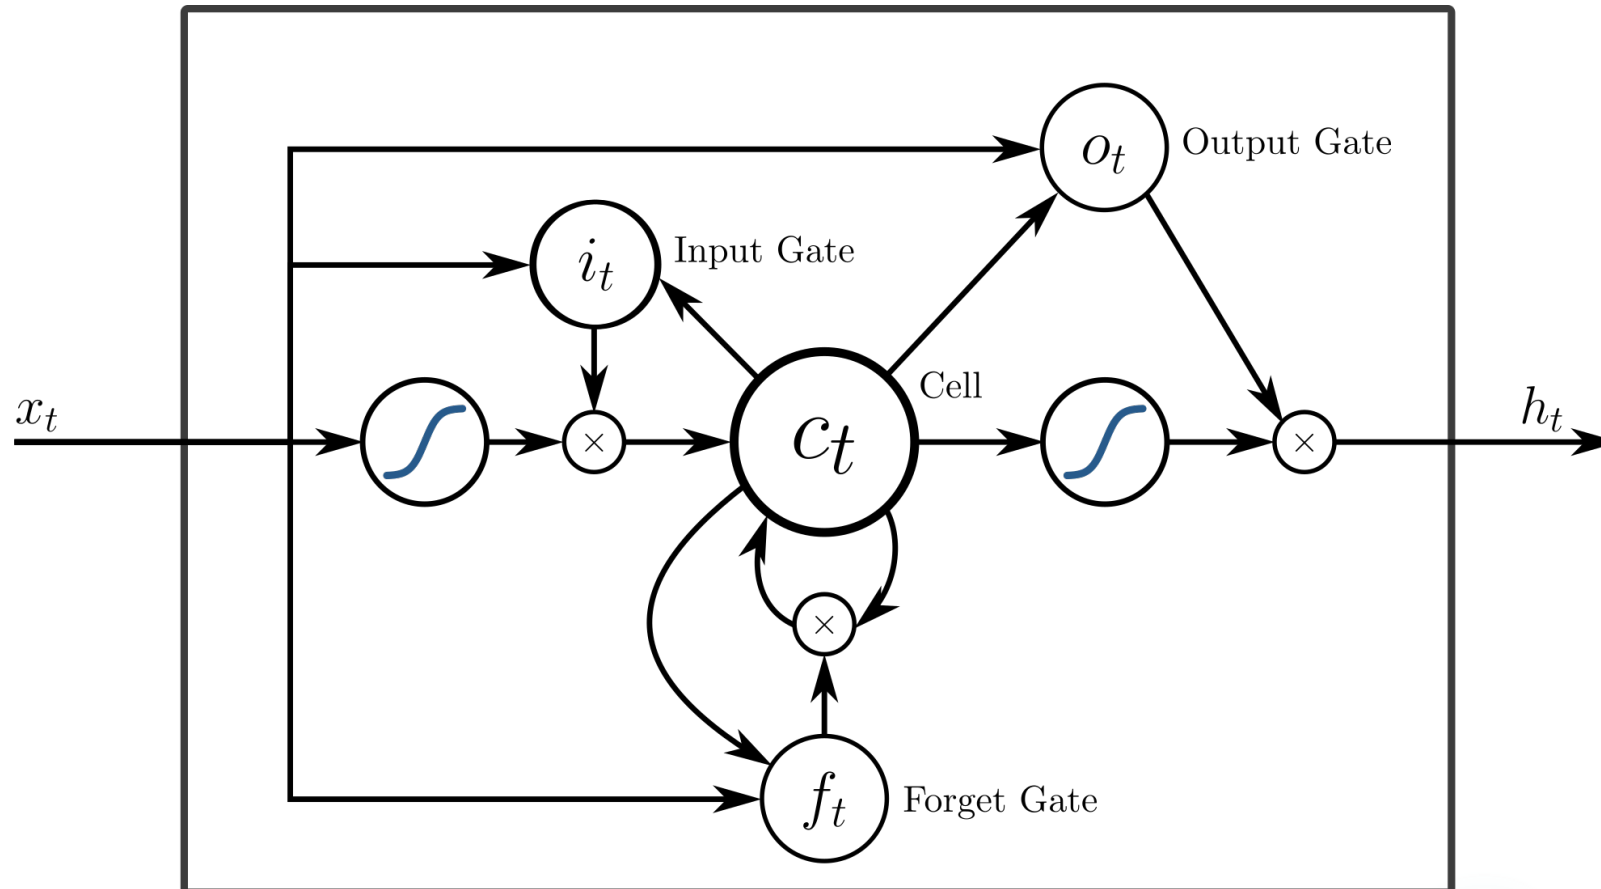


LONG SHORT-TERM MEMORY NETWORKS

- ▶ LSTMs are a special kind of RNN (invented in 1997)
- ▶ “State of the art” (the idea is old but the available computing power is new) for many sequence to sequence mapping and text generation tasks
- ▶ Adds an explicit "memory" unit
- ▶ Augment RNNs with a few additional **Gate Units**
 - ▶ Gate Units control how long/if events will stay in memory
 - ▶ **Input Gate**: If its value is such, it causes items to be stored in memory
 - ▶ **Forget Gate**: If its value is such, it causes items to be removed from memory
 - ▶ **Output Gate**: If its value is such, it causes the hidden unit to feed forward (output) in the network




LONG SHORT-TERM MEMORY NETWORKS



BIDIRECTIONAL RECURRENT NEURAL NETWORKS

- ▶ Bidirectional RNNs simply connect in both directions
- ▶ Thus, output can be dependent on both future and past inputs
- ▶ Good for context around a word, for instance
 - ▶ e.g. Named Entity Recognition, is this a "person" token?

Bidirectional RNN (BRNN)




Must learn weights w_2 , w_3 , w_4 & w_5 ; in addition to w_1 & w_6 .

Alex Graves, ["Supervised Sequence Labelling with Recurrent Neural Networks"](#)



CONVOLUTIONAL AND POOLING LAYERS

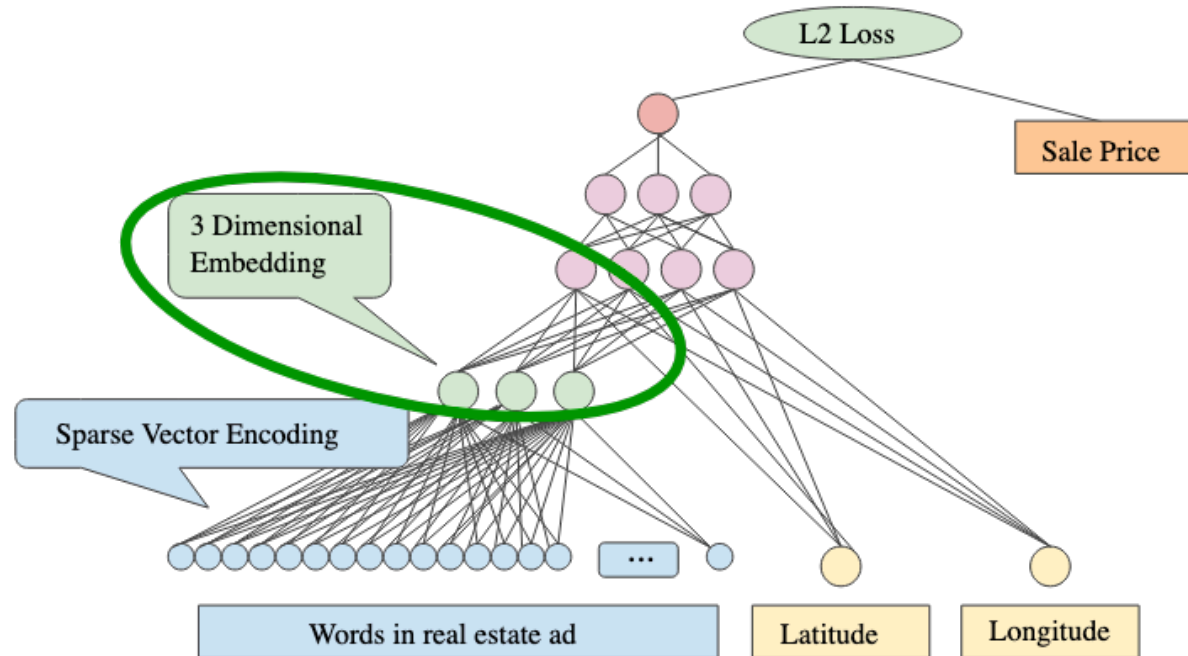
Commonly used for Image data, but they work for NLP too!



EMBEDDINGS, MERGE LAYERS

- ▶ Embeddings turn positive integers (indexes) into dense vectors of fixed size... perfect for mapping tokenized word vectors into a latent space.
- ▶ Merge layers combine outputs from multiple layers, allowing us to concatenate network branches.

(These are separate concepts but this image contains good examples of both layers.)





THANK YOU!
