
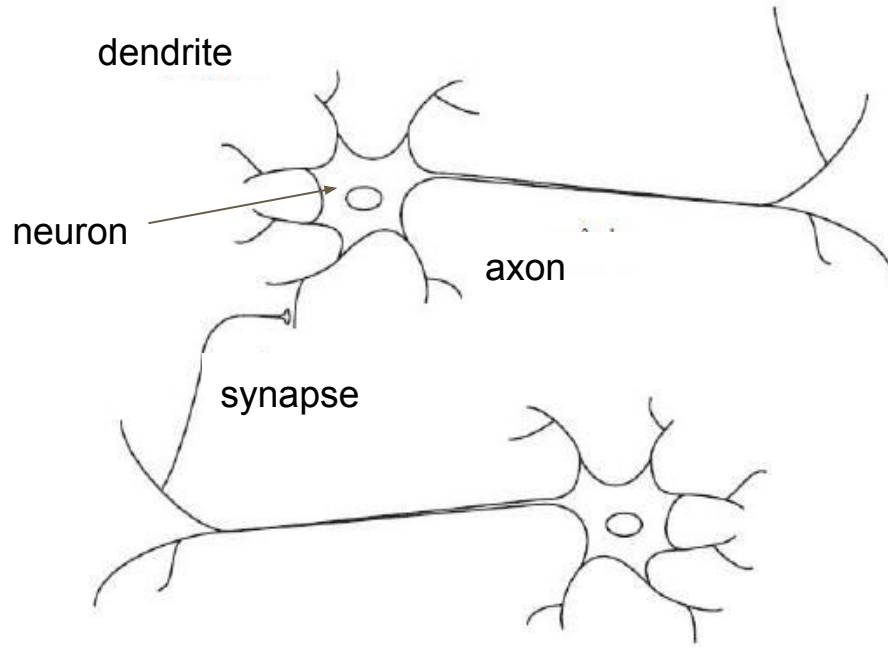


Navigating the Neural Net Terrain



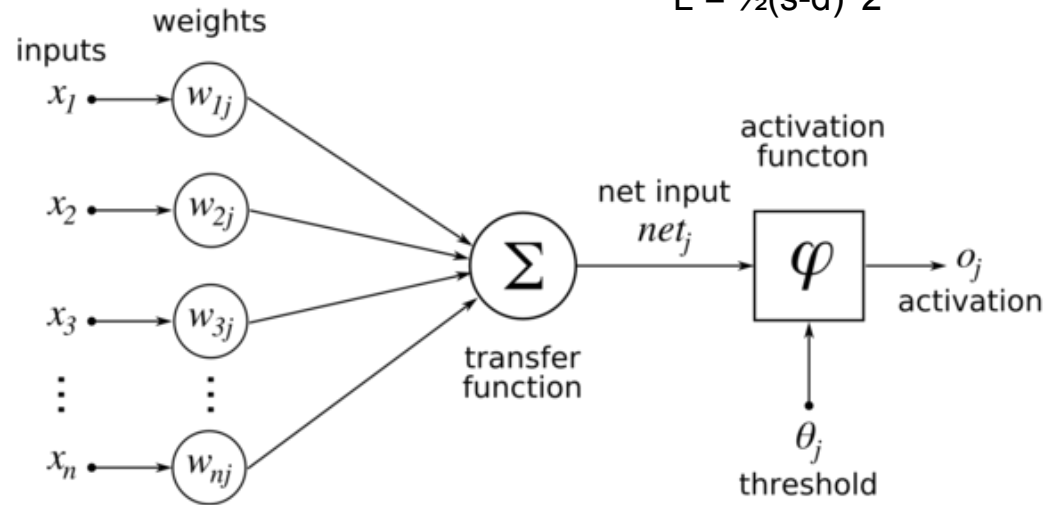
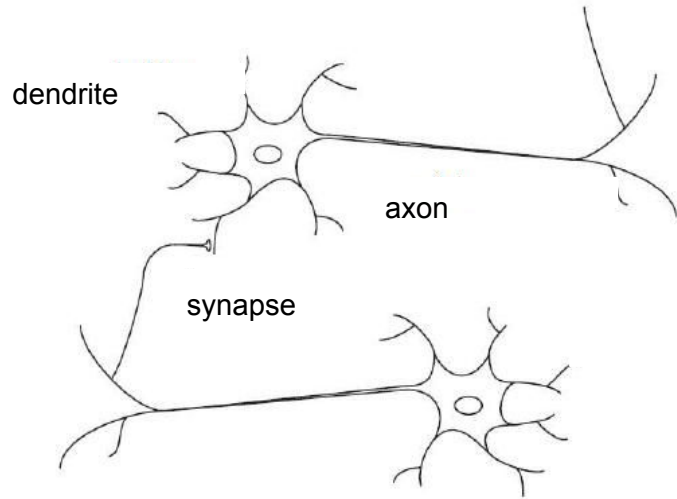
The Brain Analogy

(our cartoon neuron)



The Brain Analogy

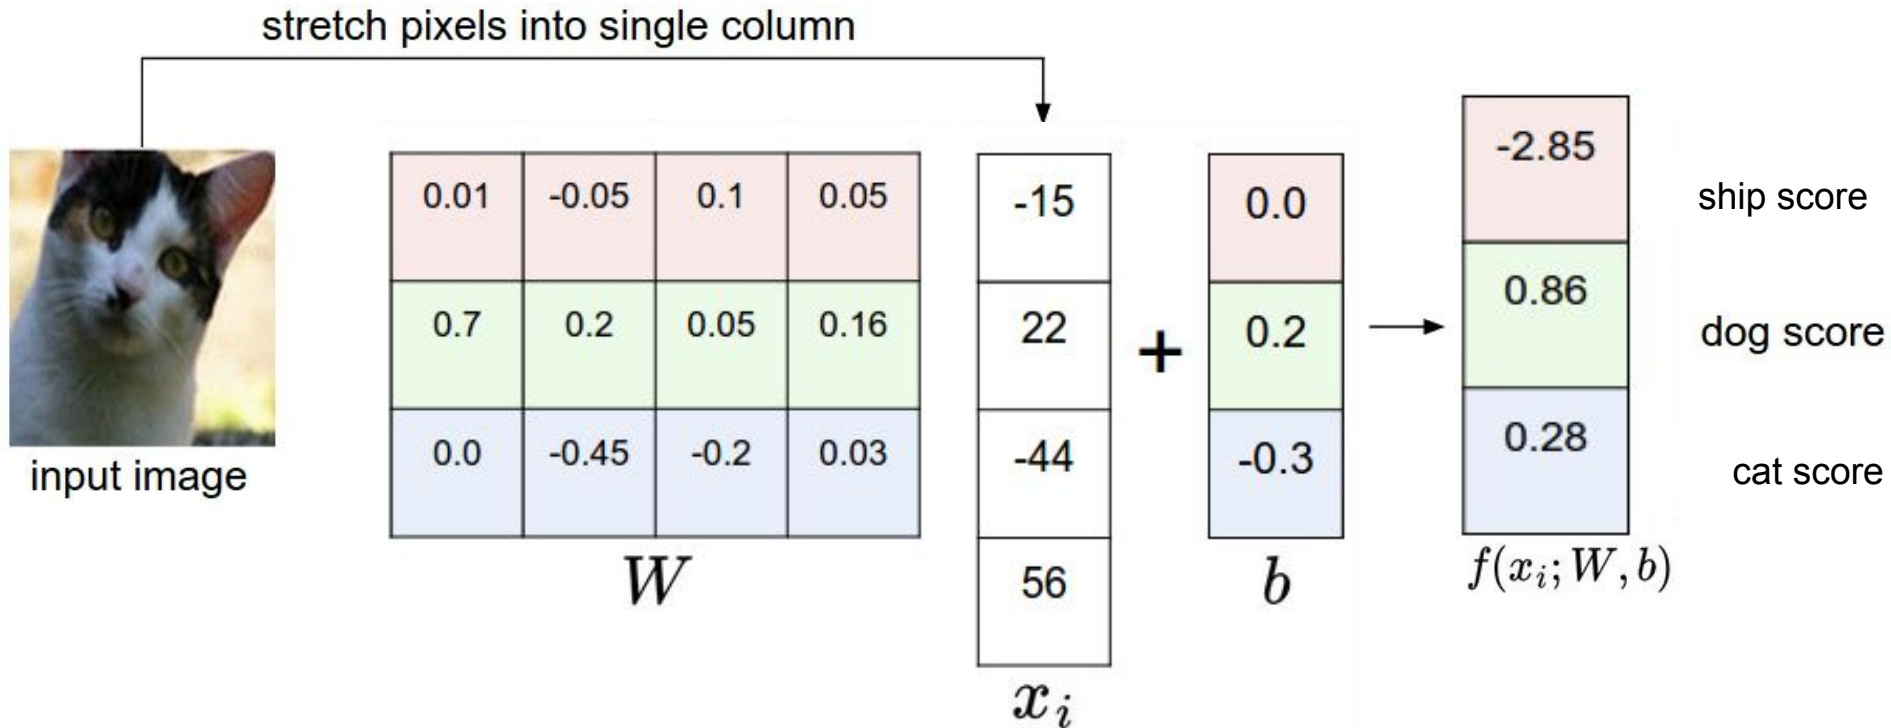
(cartoon neuron & mathematical neuron)




$$s = f(x, w)$$

$$L = \frac{1}{2}(s - d)^2$$

The Linear Classifier Analogy



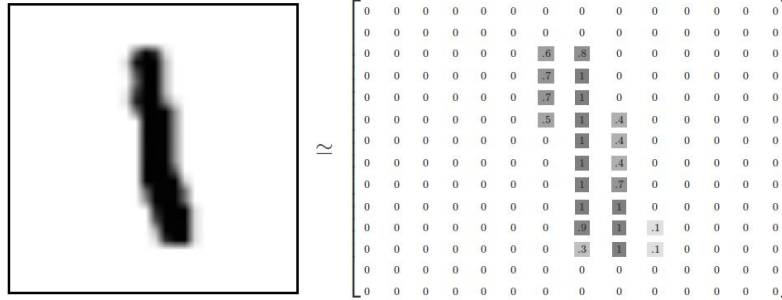
Losses: Softmax (Cross-Entropy) Loss



Softmax: $f_j(z) = \frac{e^{z_j}}{\sum_k e^{z_k}}$

Cross-Entropy $Li = -\log\left(\frac{e^{z_j}}{\sum_k e^{z_k}}\right)$


The Linear Classifier Analogy: MNIST data set



$$\begin{pmatrix} W_{1,1}x_1 + W_{1,2}x_2 + W_{1,3}x_3 + b_1 \\ W_{2,1}x_1 + W_{2,2}x_2 + W_{2,3}x_3 + b_2 \\ W_{3,1}x_1 + W_{3,2}x_2 + W_{3,3}x_3 + b_3 \end{pmatrix} = \begin{bmatrix} y_1 \\ y_2 \\ y_3 \end{bmatrix}$$

1) Add up evidence of input being in a certain class Evidence = $\sum W_{i,j}x_j + b_i$


The Linear Classifier Analogy: MNIST data set




- 1) Add up evidence of input being in a certain class Evidence = $\sum W_{i,j}x_j + b_i$

The Linear Classifier Analogy: MNIST data set


$W \approx$



0




1




2




3




4



5




6



7




8



9

BackPropagation

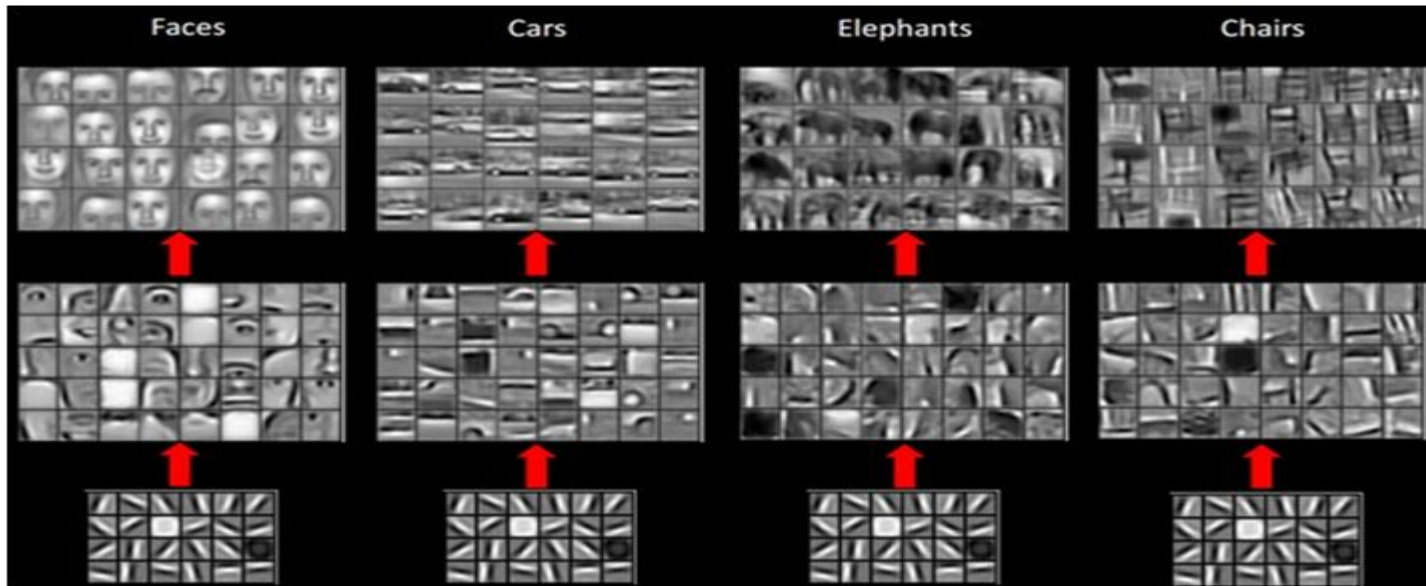
How do we get there?



Chain Rule!

Convolutional Neural Nets



Very similar to Neural Nets.. But how are they different ?



Convolutional Neural Nets

Vs. Neural Nets

- Input is an image: Leverage 3D Structure
- Fully Connected ? Not really



The CNN Family

Winners of the ILSVRC ImageNet challenges

AlexNet (2012, Krizhevsky): Popularized CNNs - 1st to incorporate consecutive convolutional layers


GoogLeNet / Inception (2014, Szegedy): Drastically reduced the # of parameters used (from 60 million to 4 million)

ResNet (2015, Kaiming He): Residual Network : famous for skip-connections and heavy use of batch-normalization; also removes some fully connected layers (at end of network)




Convolutional Neural Nets : Architecture

- 1) **Input Layer:** Raw pixel values of the image
(ex: $224 \times 224 \times 3$ (3 ~ color channels (RGB)))
- 2) Conv Layer
- 3) Pool Layer
- 4) ReLU Layer
- 5) FC (Fully Connected Layer)



Convolutional Neural Nets : Architecture


- 1) Input Layer: Raw pixel values of the image
- 2) **Conv Layer: Dot product between weights and the small region of input volume (ex: 11 x 11 x 3 filters)**
- 3) Pool Layer
- 4) ReLU Layer
- 5) FC (Fully Connected Layer)



Convolutional Neural Nets : Architecture


What is a Convolution?

$$f * g = \int f(t - \tau)g(\tau)d\tau$$



Convolutional Neural Nets : Architecture

What is a Convolution?



1 1 1 0
1 1 1 0
1 1 1 0

1 1 1
1 -8 1
1 1 1

→ 0

1 1 1 0
1 1 1 0
1 1 1 0

1 1 1
1 -8 1
1 1 1

→ -3

Convolutional Neural Nets : Architecture


What is a Convolution?

Convolutional Layer: $(W-F + 2P)/S + 1$

- **W** : Input Volume size
- **F**: Receptive Field size of the Conv Layer Neuron
- **P**: Zero- Padding
- **S**: Stride

$$(224 - 11 + 2(3))/4 + 1 = 55$$

Conv Layer Output ~ 55 x 55 x 96 (ie : 55² neurons in each layer)



Convolutional Neural Nets : Architecture

What is a Convolution?

Voila. We have 96 filters.




Convolutional Neural Nets : Architecture

- 1) Input Layer
- 2) Conv Layer
- 3) **Pooling Layer: Performs downsampling operation**
- 4) ReLU Layer
- 5) FC (Fully Connected Layer)

Our Eqn : $O = (W - F) / S + 1$


AlexNet: use 3 x 3 MaxPooling w/ stride = 2

$$O = (55 - 3) / 2 + 1 = 27$$



Convolutional Neural Nets : Architecture

- 1) Input Layer: Raw pixel values of the image
- 2) Conv Layer:
- 3) Pool Layer:
- 4) **ReLU Layer: Apply an elementwise activation function**
(ex : $\max(0, x)$ thresholding output dimension ~ same as input)
- 5) FC (Fully Connected Layer)



*The ReLU non-linearity is applied to the output of every convolutional and fully-connected layer.


Convolutional Neural Nets : Architecture

ReLU Layer:

Traditionally:

$f(x) = \tanh(x)$ or $fx = (1+e^{-x})^{-1}$ (Very slow to train)


$f(x) = \tanh(x)$



Now:


$f(x) = \max(0, x)$ (Faster to train)

$f(x) = \max(0, x)$

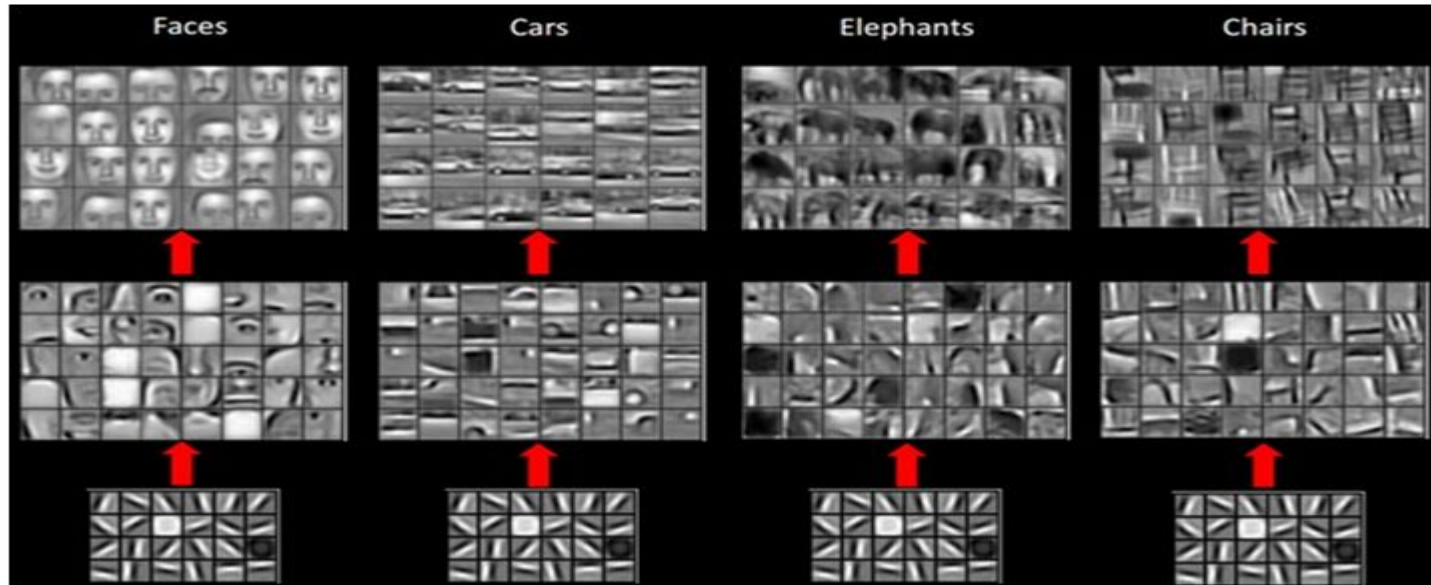


Convolutional Neural Nets : Architecture

- 1) Input Layer: Raw pixel values of the image
- 2) Conv Layer:
- 3) Pool Layer:
- 4) ReLU Layer:
- 5) **FC (Fully Connected) Layer** : Each neuron will be connected to all activations of the previous volume. The output layer will compute class scores (ex: $[1 \times 1 \times 1000]$)



Convolutional Neural Nets : Architecture




Real Life example #1:

Train Retinopathy data via AlexNet

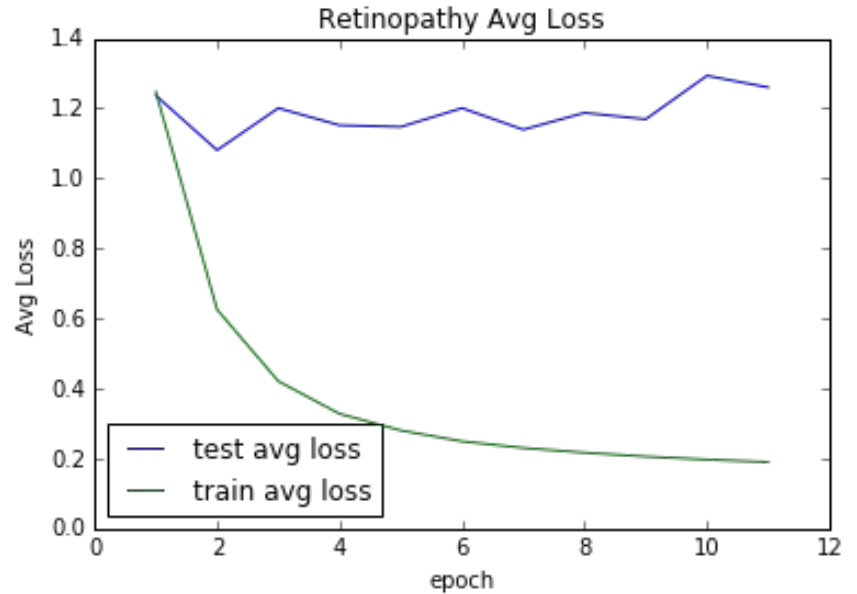

Kaggle Dataset:

Diabetic Retinopathy Detection



Learning from the Learning Process

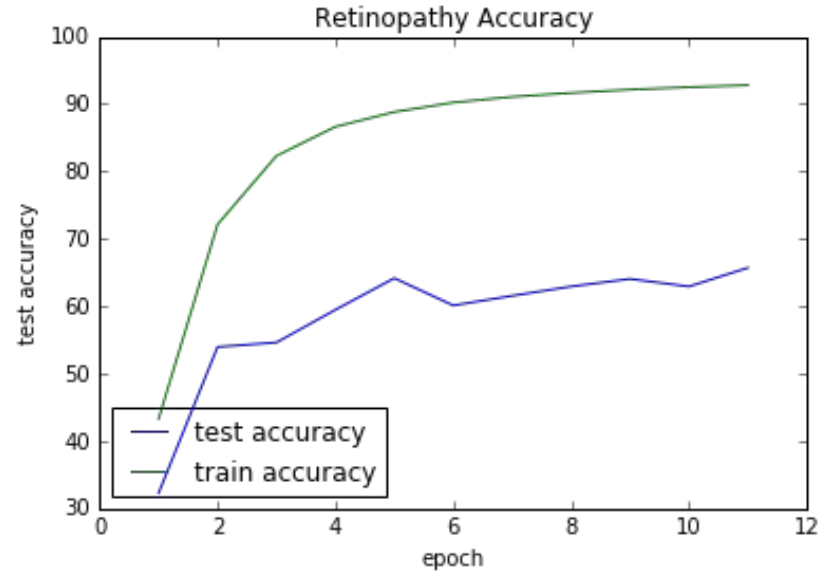

1) Loss functions



* Tip: Change the learning rate!

Learning from the Learning Process

2) Accuracy



* Tip: Increase L2 weight penalty , Increase Drop-Out, More Data (possibly with jitter) - -try batch norm ?

Learning from the Learning Process:

3) Weight Ratios


- update / weight ratio : should be roughly about $1e-3$

(larger ratio~ learning rate may be too high, lower ratio~ learning rate may be too low)

4) First-layer Visualizations

Visualized weights from the 1st layer of the network:


(smooth, diverse features indicate that training is going well)



Real Life example #2:


Can I leverage the Inception model to decipher different lung diseases?

Train XRAY data via Inception-V3




Opacity

Train XRAY data via TF's Inception-V3



Train using Inception-V3 from scratch
(not looking good)



Retrain using Inception-V3 from pretrained model
(much better!)

Why did that work so well ???

